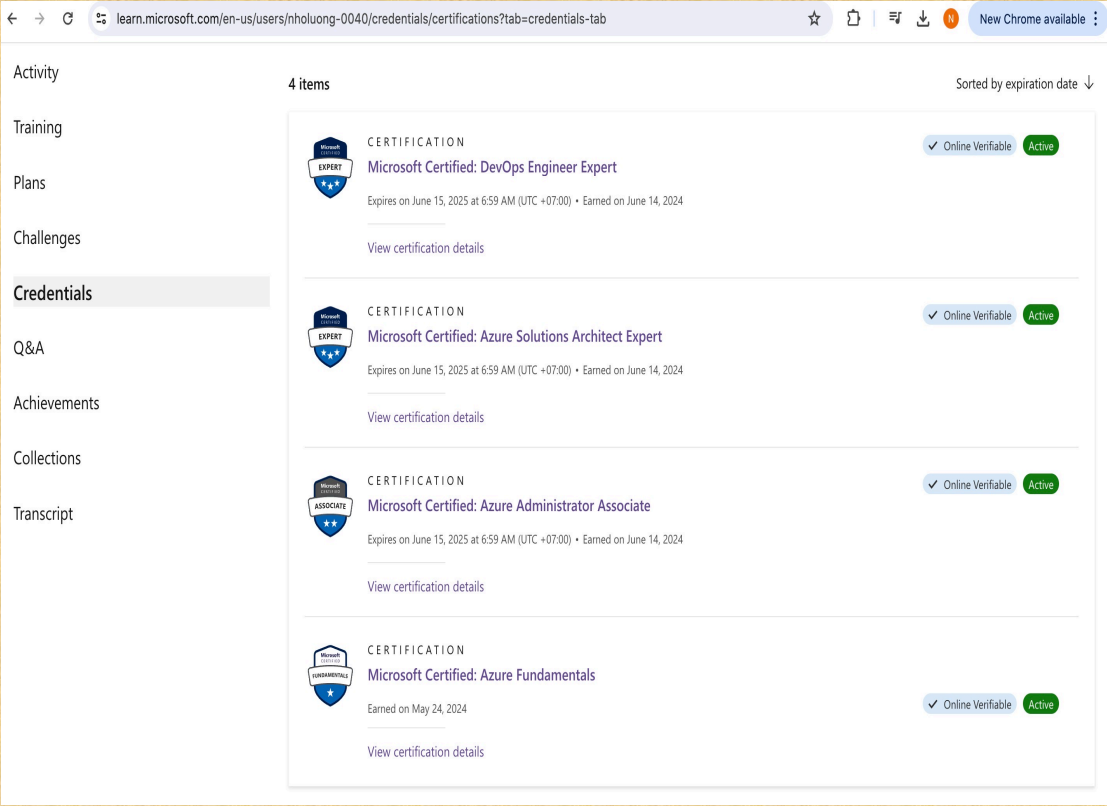
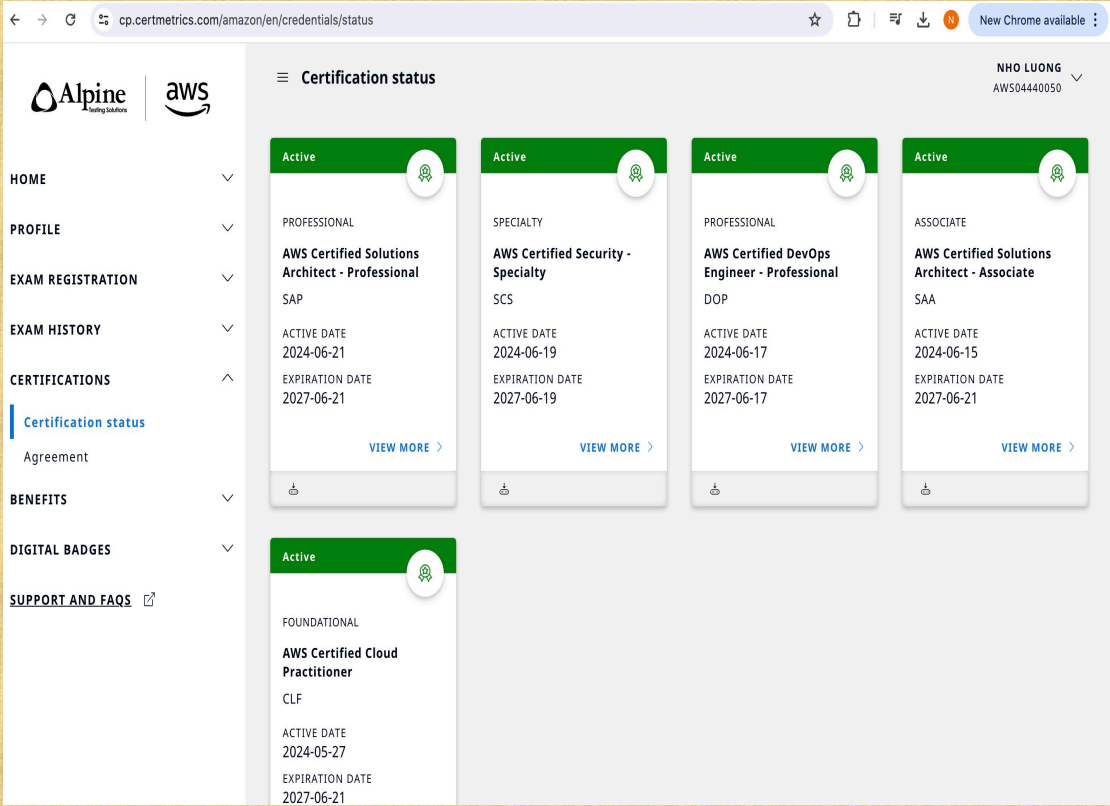


# Kubernetes for Beginners on AWS Cloud

Author: Nho Luong

Skill: DevOps Engineer Lead



Author: Nho Luong

Skill: DevOps Engineer Lead

# Docker Architecture





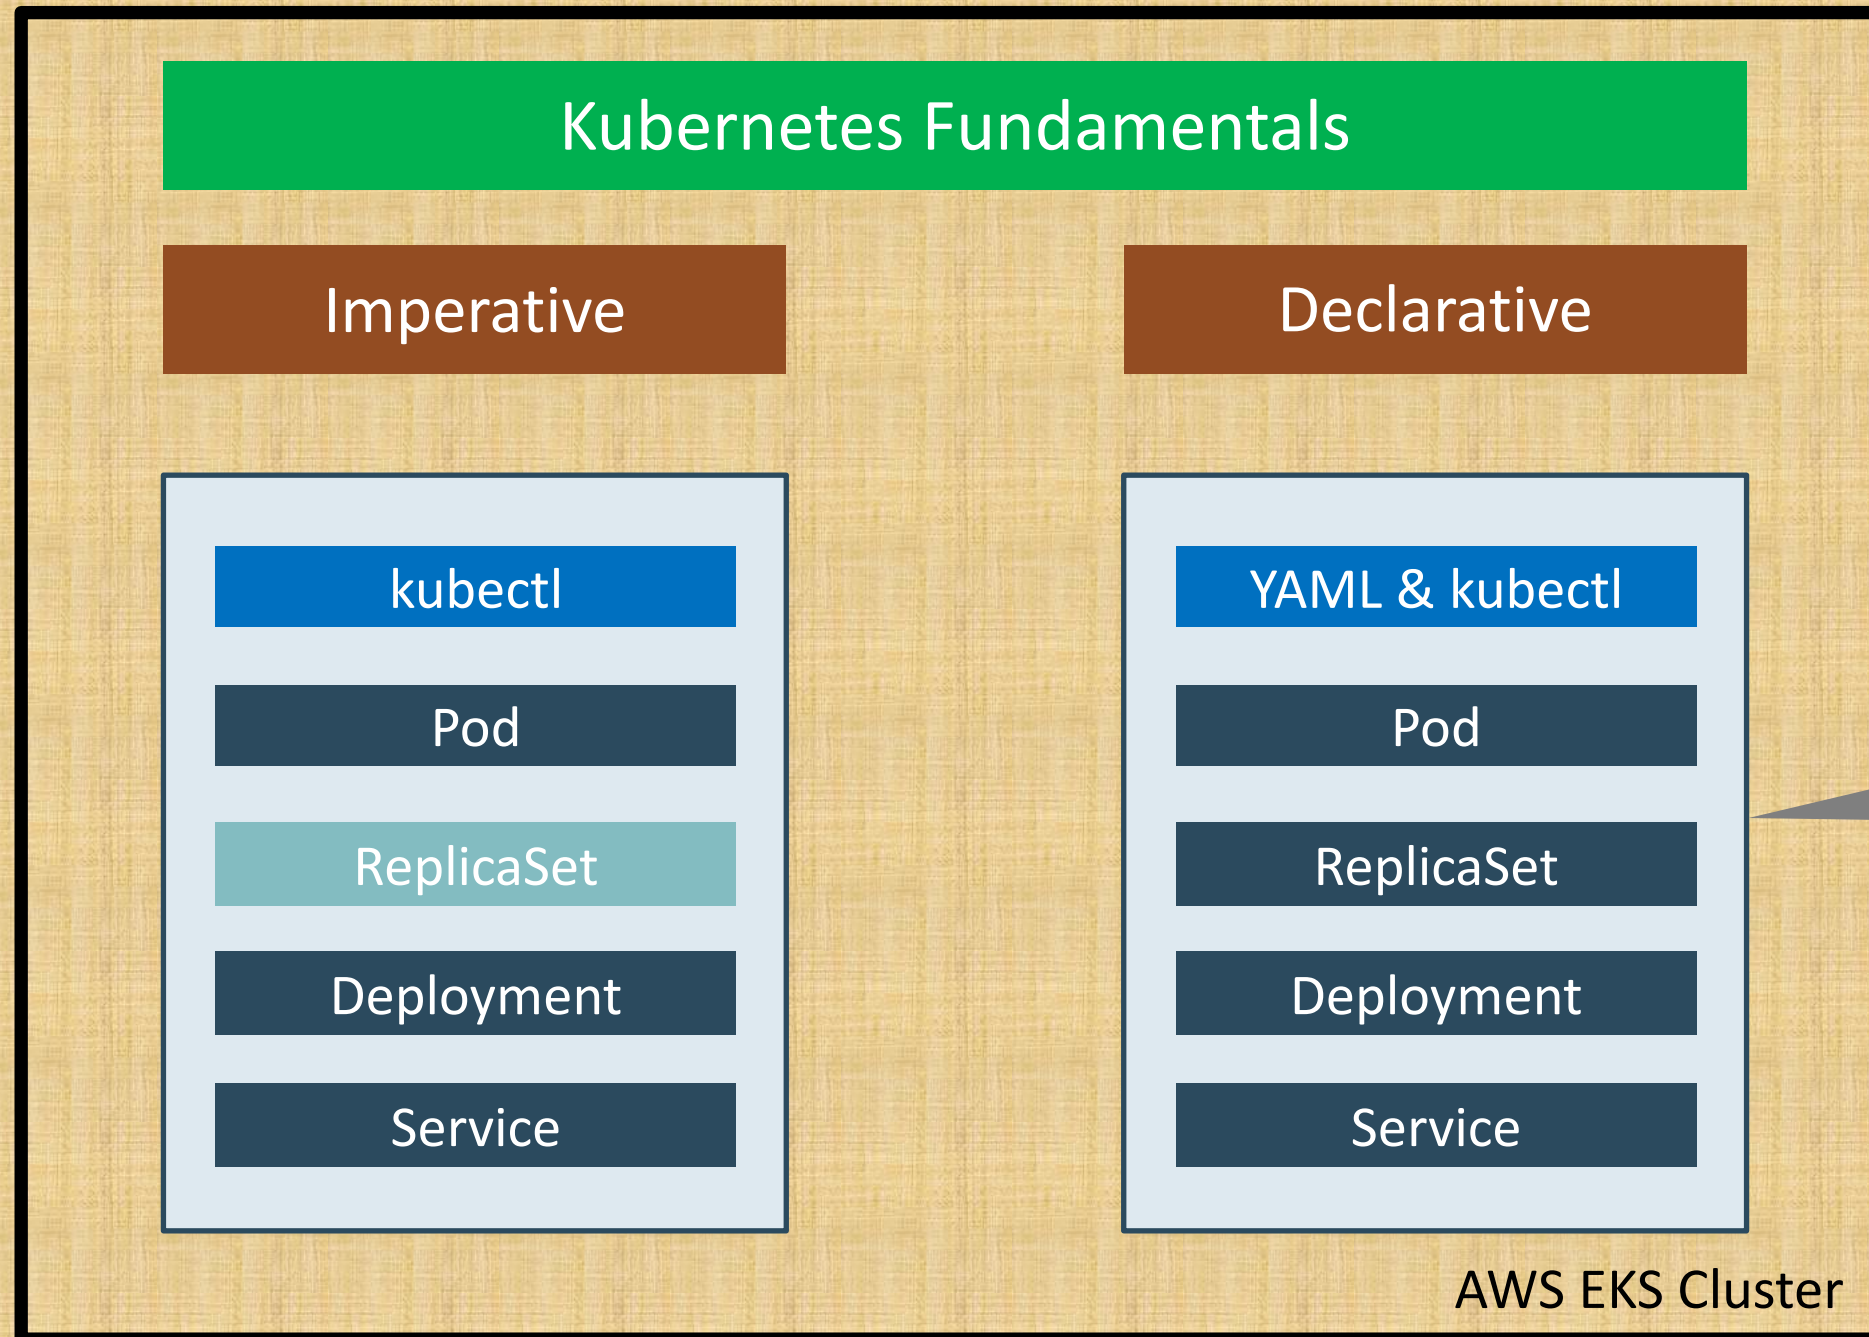
# Kubernetes on AWS Cloud Course Outline



Author: Nho Luong

Skill: DevOps Engineer Lead

# Kubernetes on AWS Cloud





Free Courses – 2 Hours limitation on Udemy

Kubernetes for Absolute Beginners on AWS Cloud | Part-1

Kubernetes for Absolute Beginners on AWS Cloud | Part-2



# AWS EKS

## CLIs



# AWS EKS Cluster - CLIs



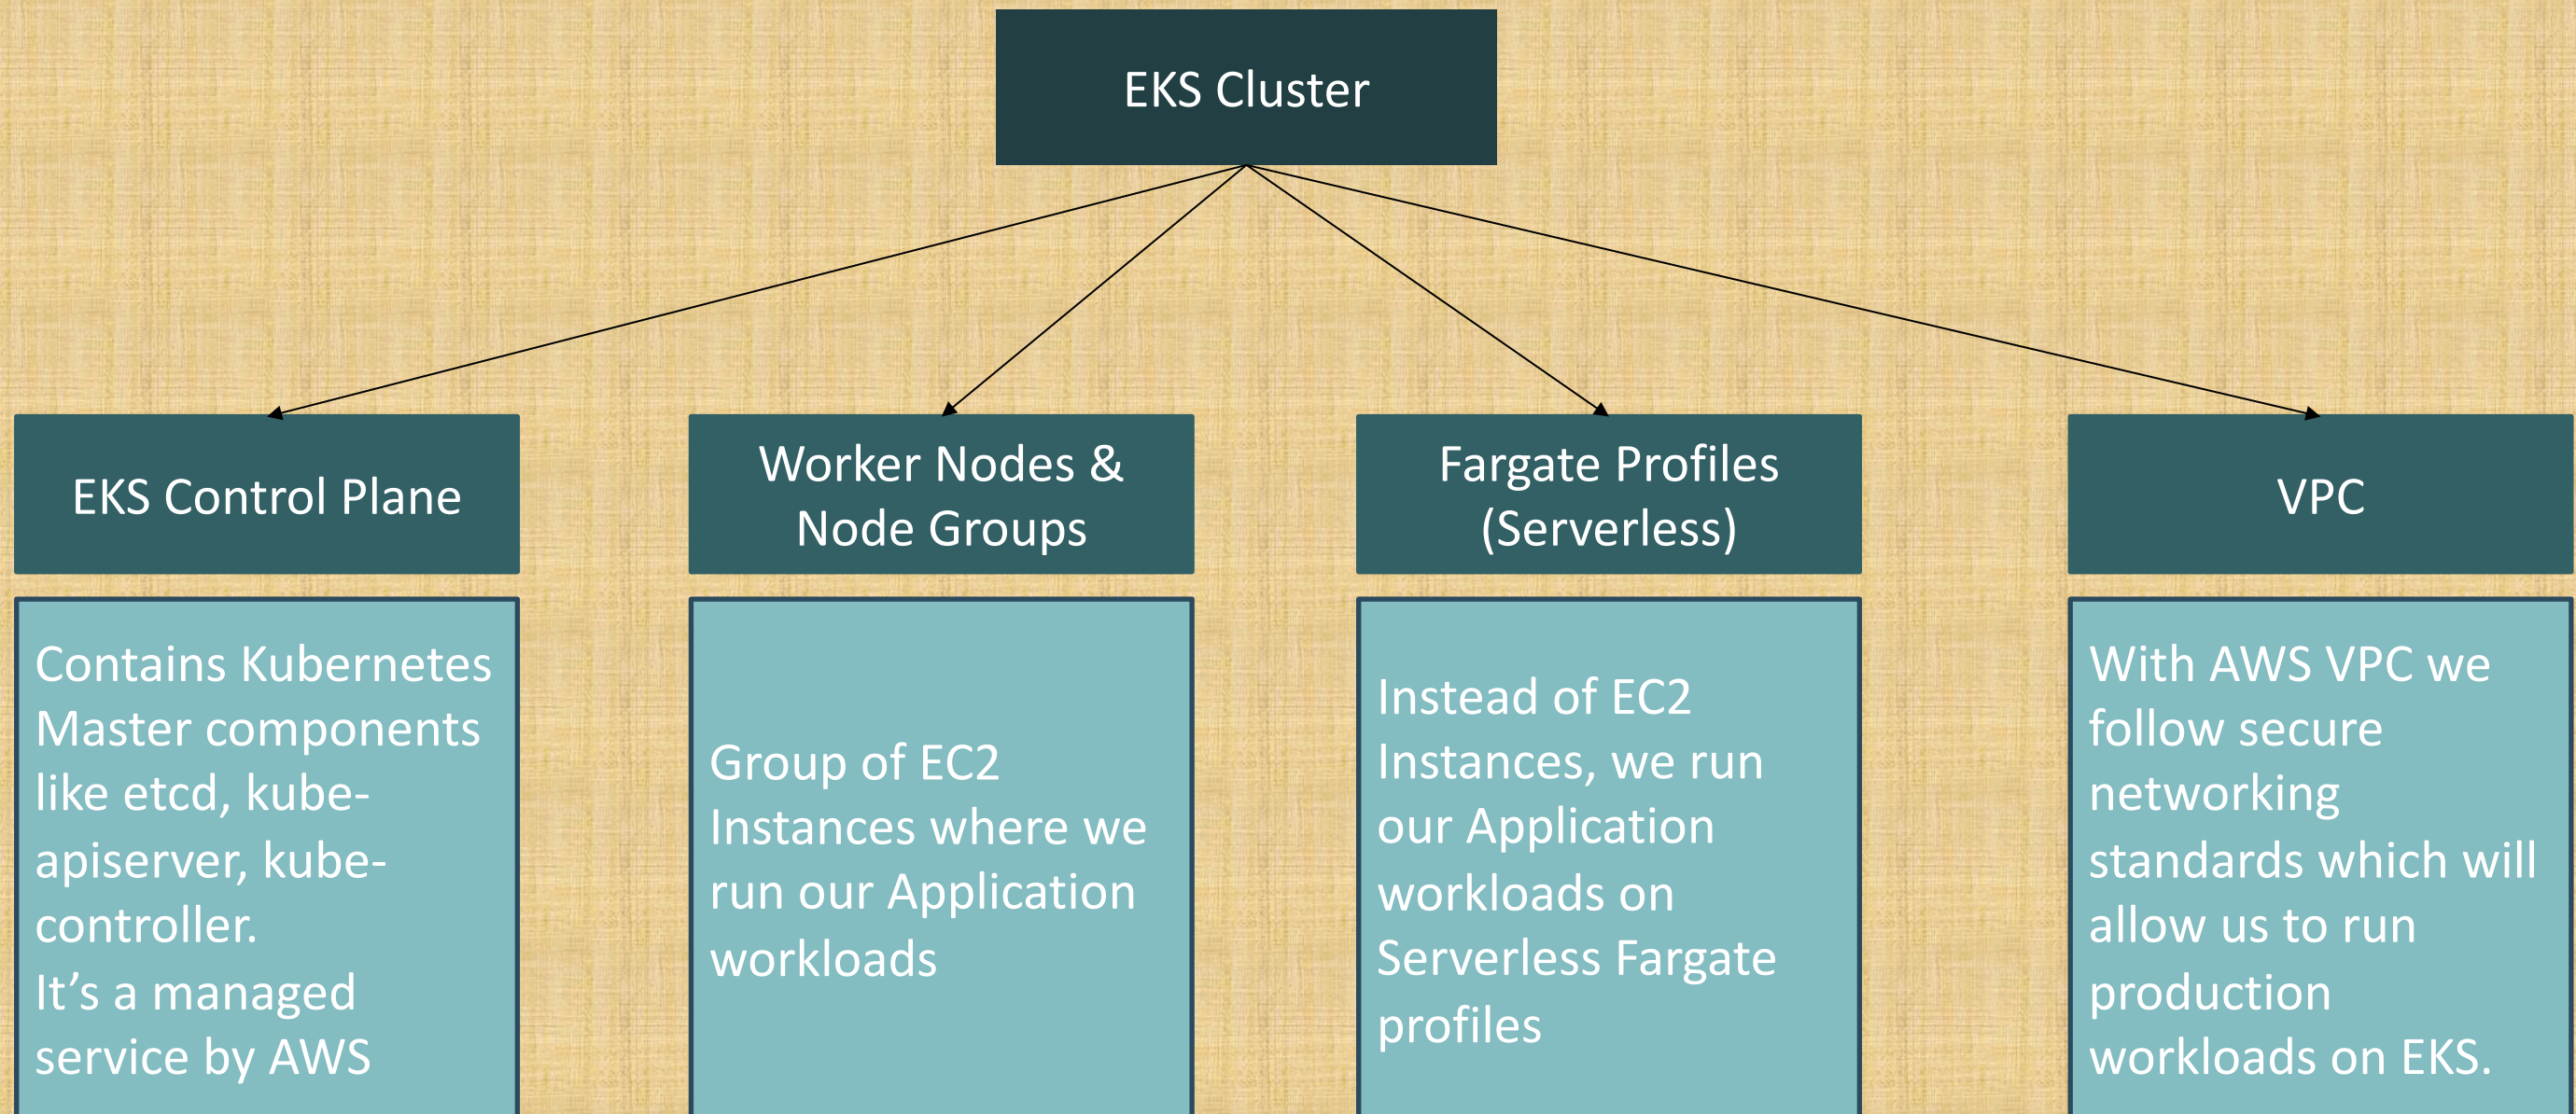


# AWS EKS Cluster

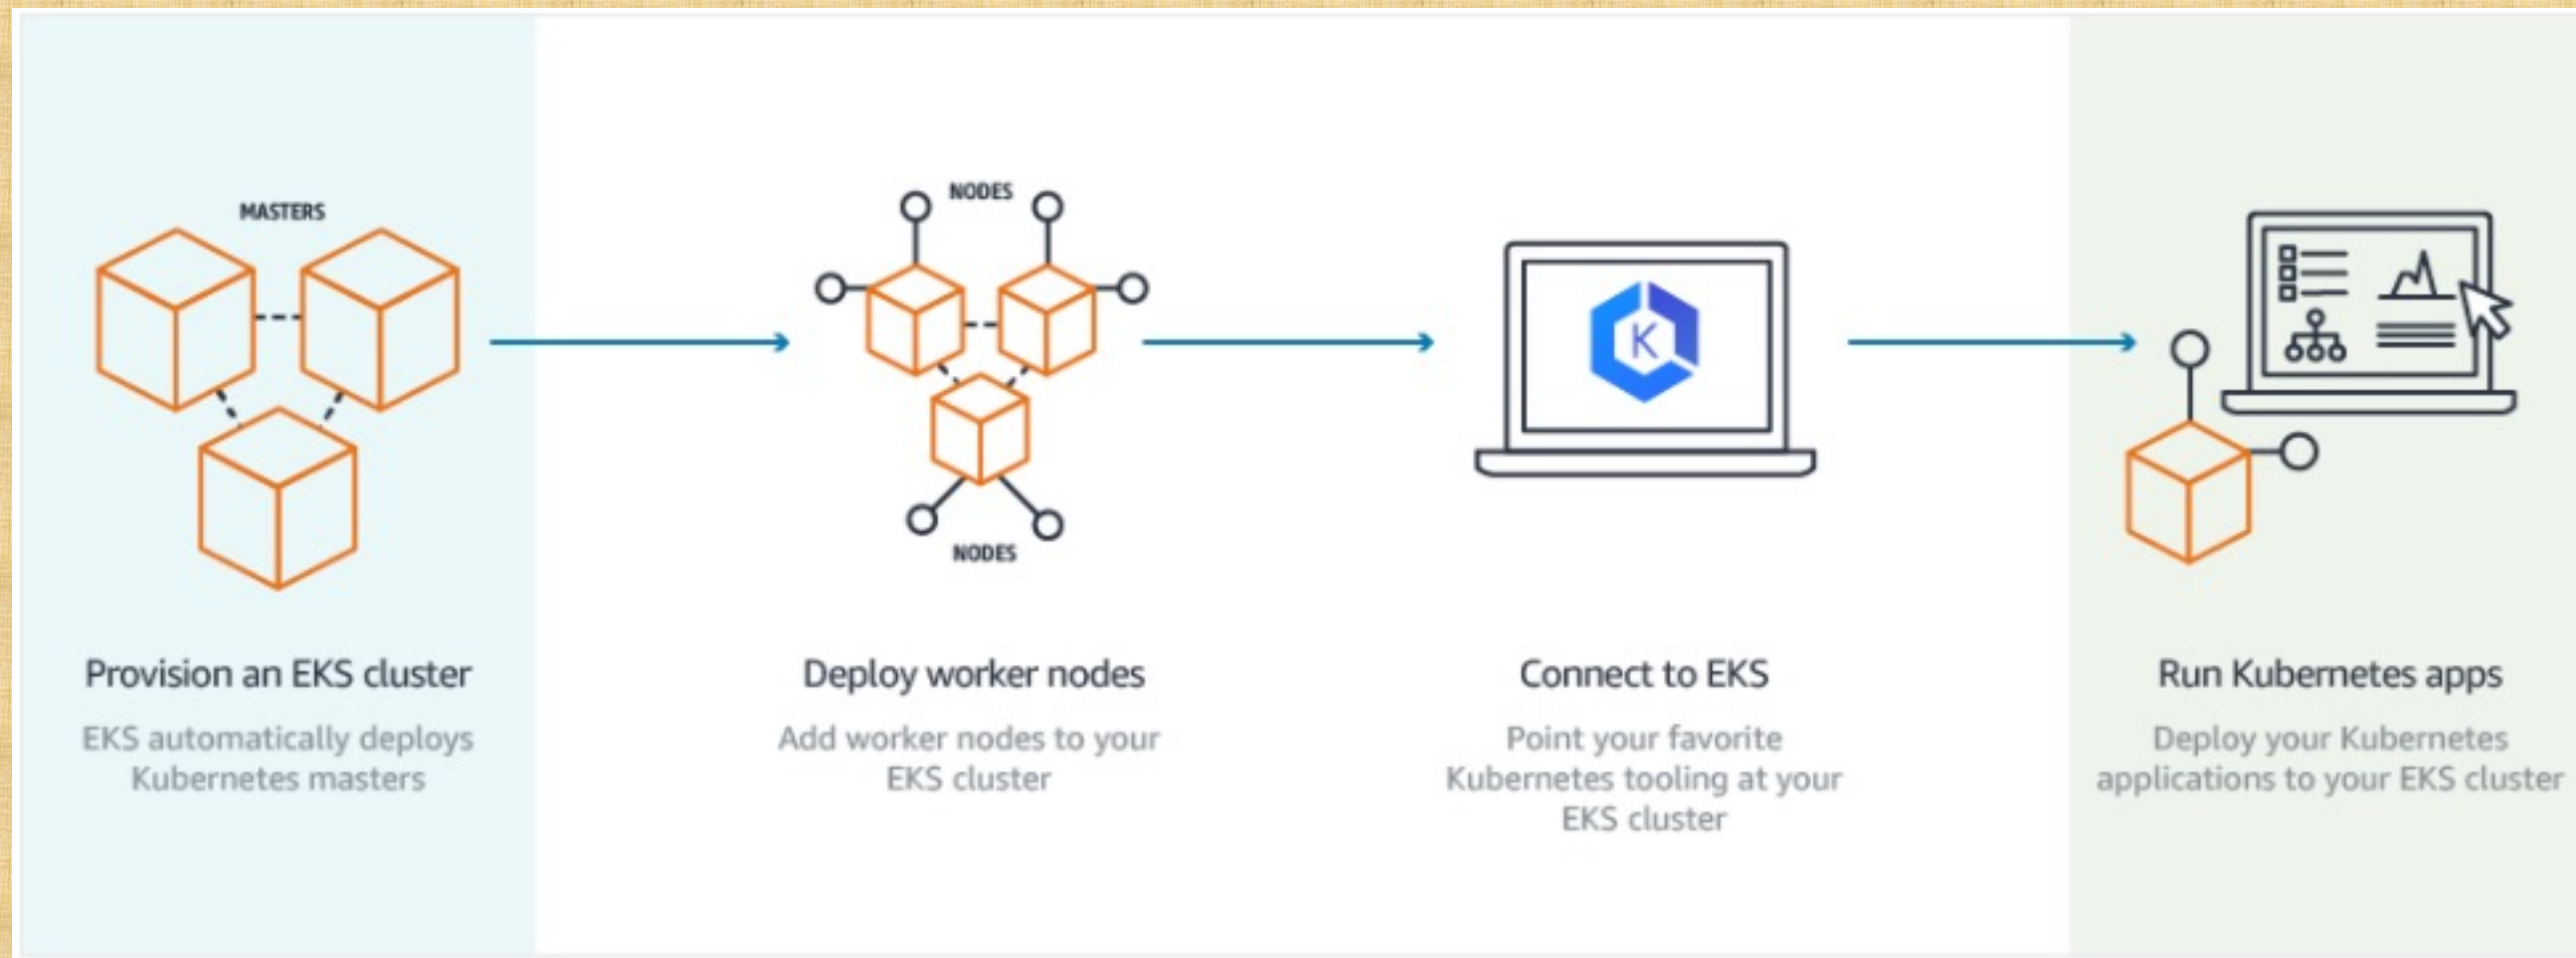




# AWS EKS – Core Objects



# How does EKS work?



# EKS Cluster – Core Objects Detailed

## EKS Control Plane

1. EKS runs a single tenant Kubernetes control plane for each cluster, and control plane infrastructure is **not shared** across clusters or AWS accounts.
2. This control plane consists of at least two API server nodes and three etcd nodes that run across **three Availability Zones within a Region**
3. EKS **automatically detects and replaces unhealthy** control plane instances, restarting them across the Availability Zones within the Region as needed.

## Worker Nodes & Node Groups

1. Worker machines in Kubernetes are called nodes. These are EC2 Instances
2. EKS worker nodes run in our AWS account and connect to our cluster's control plane via the **cluster API server endpoint**.
3. A node group is **one or more EC2 instances** that are deployed in an EC2 Autoscaling group.
4. All instances in a node group must
  1. Be the **same instance type**
  2. Be **running the same AMI**
  3. Use the **same EKS worker node IAM role**

# EKS Cluster – Core Objects Detailed

## Fargate Profiles

1. AWS Fargate is a technology that provides **on-demand, right-sized compute capacity** for containers
2. With Fargate, we **no longer** have to provision, configure, or scale groups of virtual machines to run containers.
3. Each pod running on Fargate has its **own isolation boundary** and does not share the underlying kernel, CPU resources, memory resources, or elastic network interface with another pod.
4. AWS specially built **Fargate controllers** that recognizes the pods belonging to fargate and schedules them on Fargate profiles.
5. We will see more in our Fargate learning section.

## VPC

1. EKS uses AWS VPC network policies **to restrict traffic** between control plane components to within a single cluster.
2. Control plane components for a EKS cluster **cannot view or receive** communication from other clusters or other AWS accounts, except as authorized with Kubernetes RBAC policies.
3. This **secure and highly-available configuration** makes EKS reliable and recommended for **production workloads**.



# Kubernetes Architecture







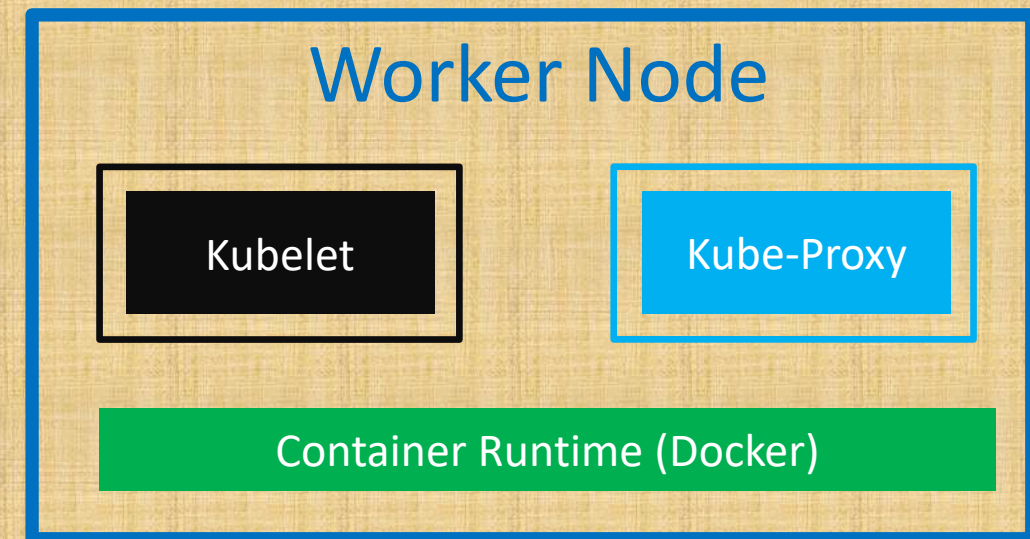
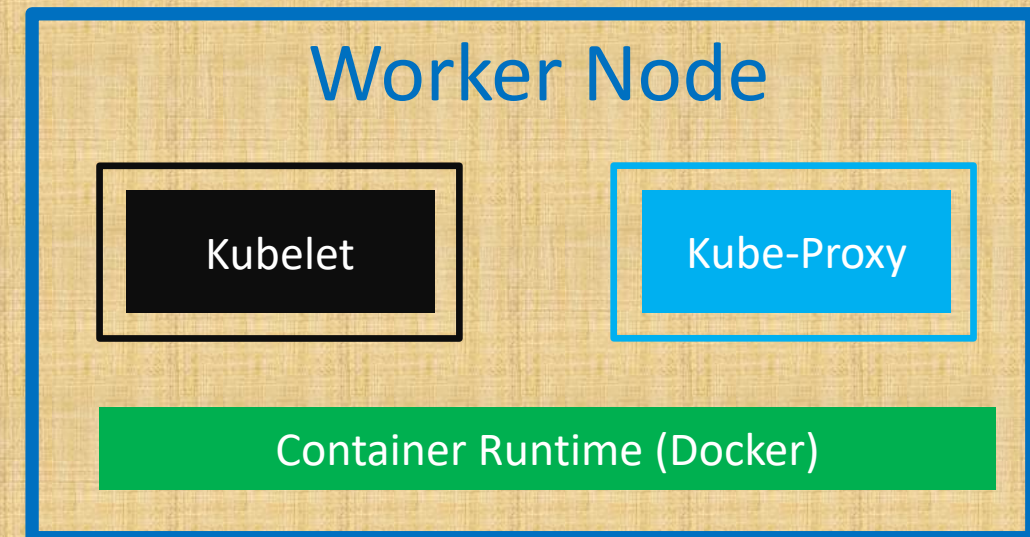
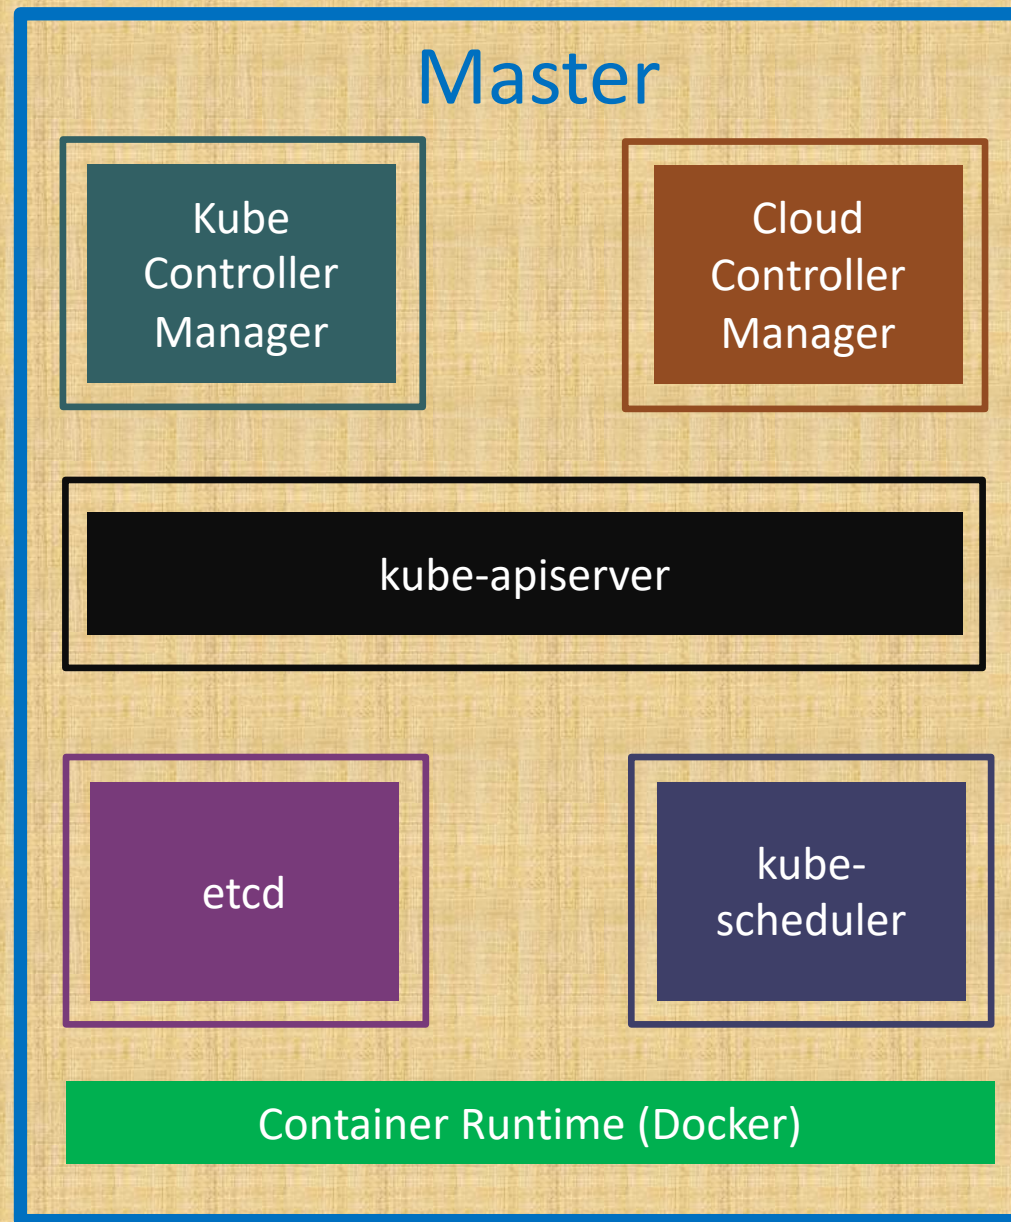
# Kubernetes Architecture



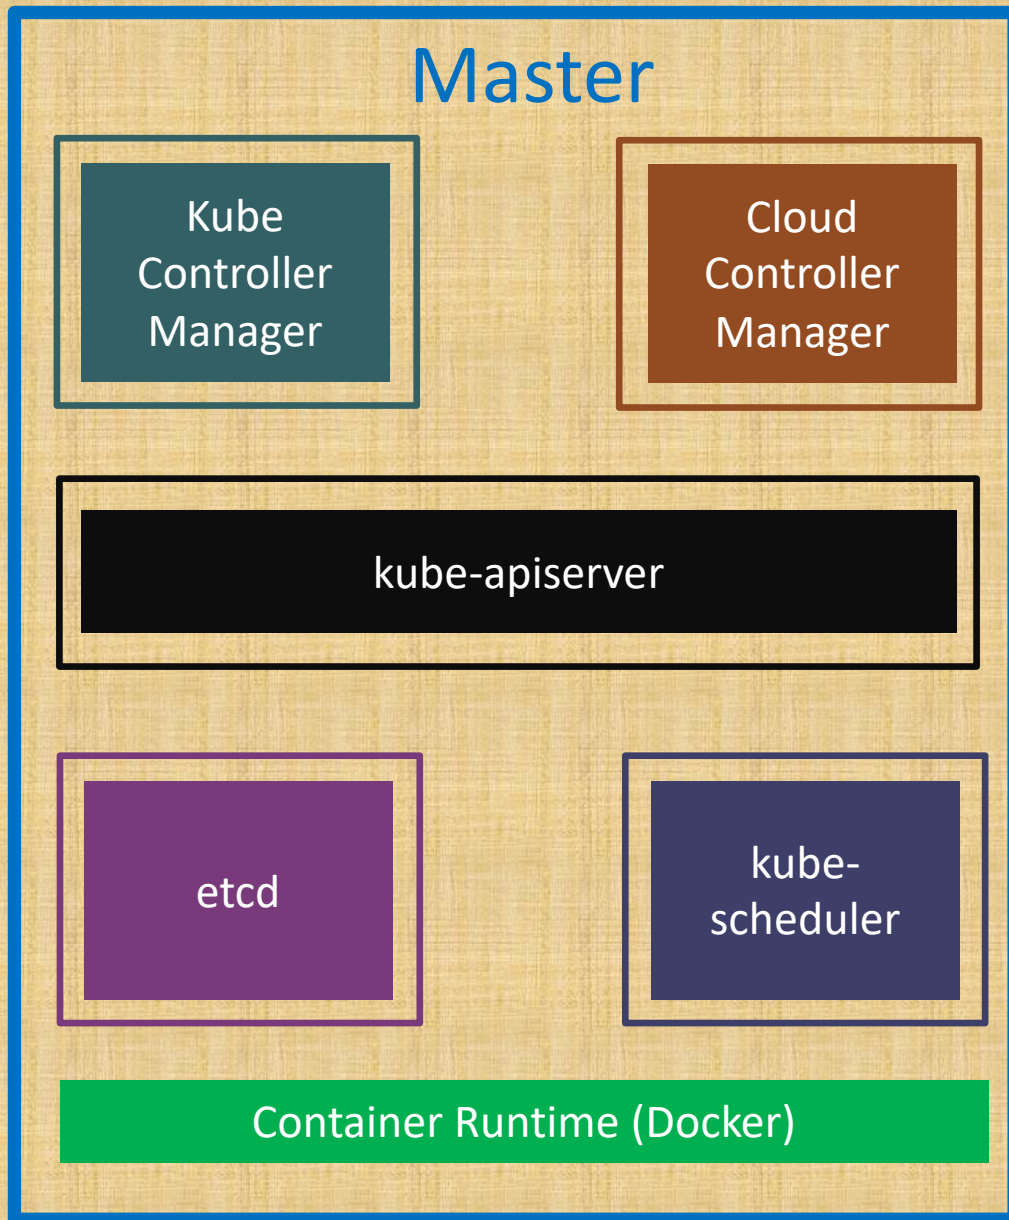
Author: Nho Luong

Skill: DevOps Engineer Lead

# Kubernetes - Architecture




# Kubernetes Architecture - Master



- **kube-apiserver**
  - It acts as **front end** for the Kubernetes control plane. It **exposes** the Kubernetes API
  - Command line tools (like kubectl), Users and even Master components (scheduler, controller manager, etcd) and Worker node components like (Kubelet) **everything talk** with API Server.
- **etcd**
  - Consistent and highly-available **key value store** used as Kubernetes' **backing store** for all cluster data.
  - It **stores** all the masters and worker node information.
- **kube-scheduler**
  - Scheduler is responsible for distributing containers across multiple nodes.
  - It watches for newly created Pods with no assigned node, and selects a node for them to run on.




# Kubernetes Architecture - Master



- **kube-controller-manager**
  - Controllers are responsible for noticing and responding when nodes, containers or endpoints go down. They make decisions to bring up new containers in such cases.
  - **Node Controller**: Responsible for noticing and responding when **nodes go down**.
  - **Replication Controller**: Responsible for maintaining the **correct number of pods** for every replication controller object in the system.
  - **Endpoints Controller**: **Populates** the Endpoints object (that is, joins Services & Pods)
  - **Service Account & Token Controller**: Creates default accounts and API Access for **new namespaces**.


# Kubernetes Architecture - Master



- **cloud-controller-manager**
  - A Kubernetes control plane component that embeds **cloud-specific control logic**.
  - It only runs controllers that are **specific** to your cloud provider.
- **On-Premise** Kubernetes clusters will not have this component.
- **Node controller**: For **checking** the cloud provider to determine if a node has been deleted in the cloud after it stops responding
- **Route controller**: For setting up **routes** in the underlying cloud infrastructure
- **Service controller**: For creating, updating and deleting cloud provider **load balancer**



# Kubernetes Architecture – Worker Nodes



- Container Runtime

- Container Runtime is the **underlying software** where we run all these Kubernetes components.
- We are using Docker, but we have other runtime options like rkt, container-d etc.

- Kubelet

- Kubelet is the **agent** that runs on every node in the cluster
- This agent is **responsible** for making sure that containers are running in a Pod on a node.

- Kube-Proxy

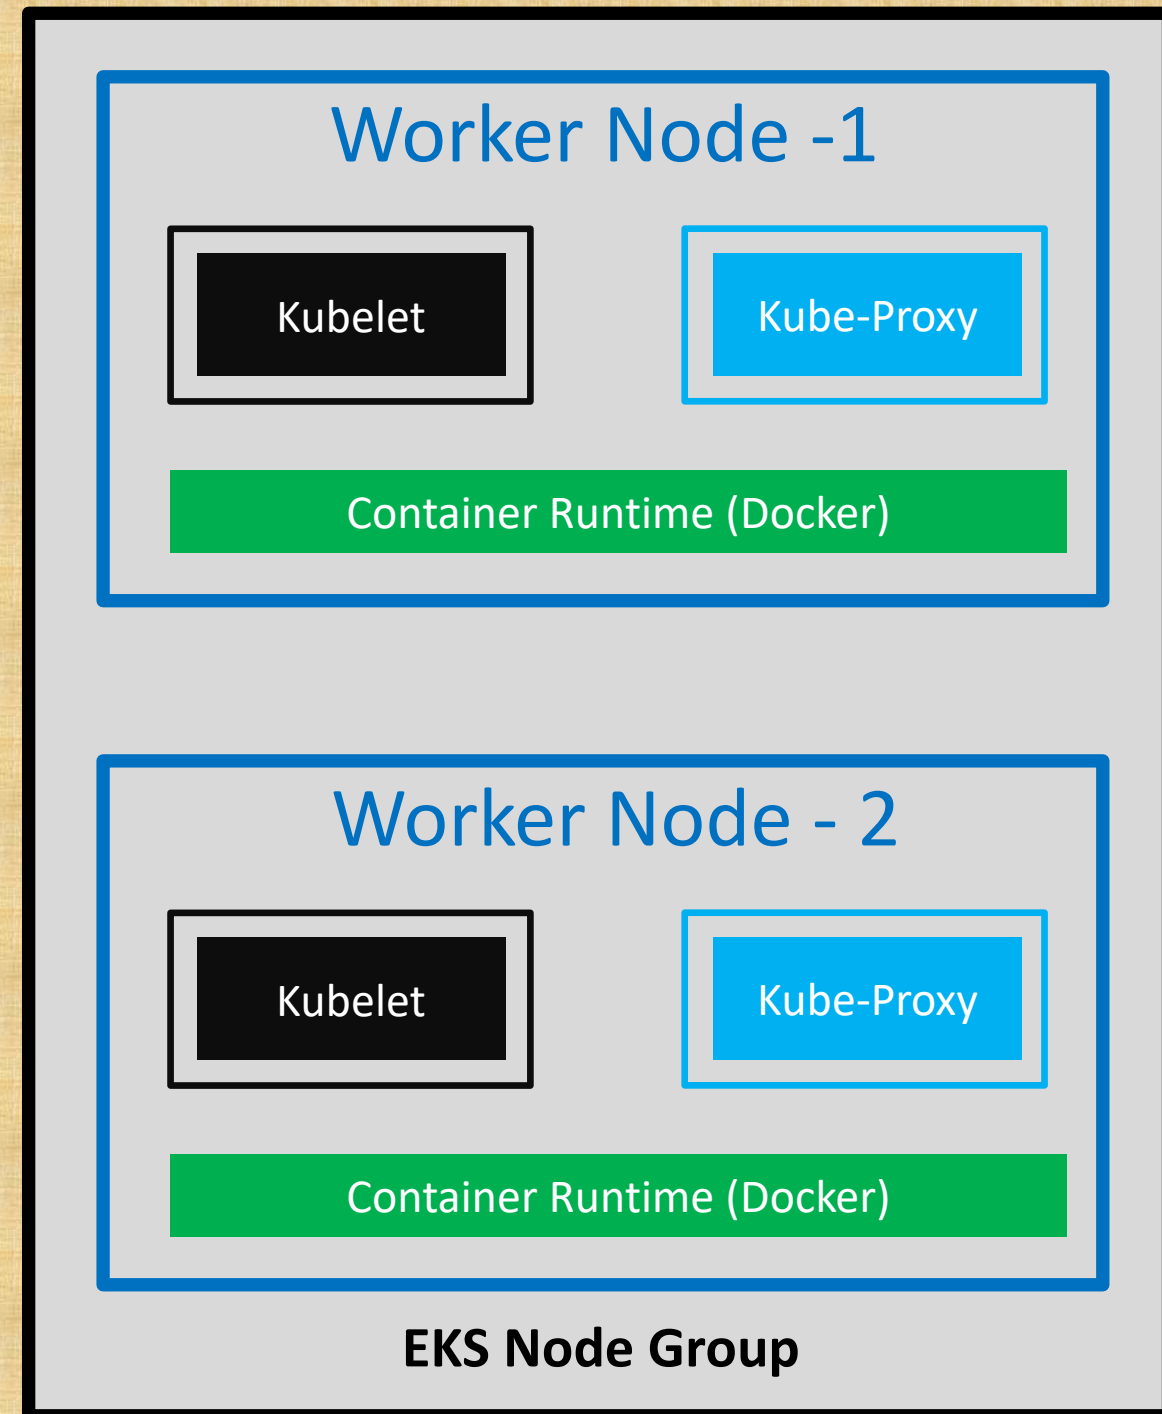
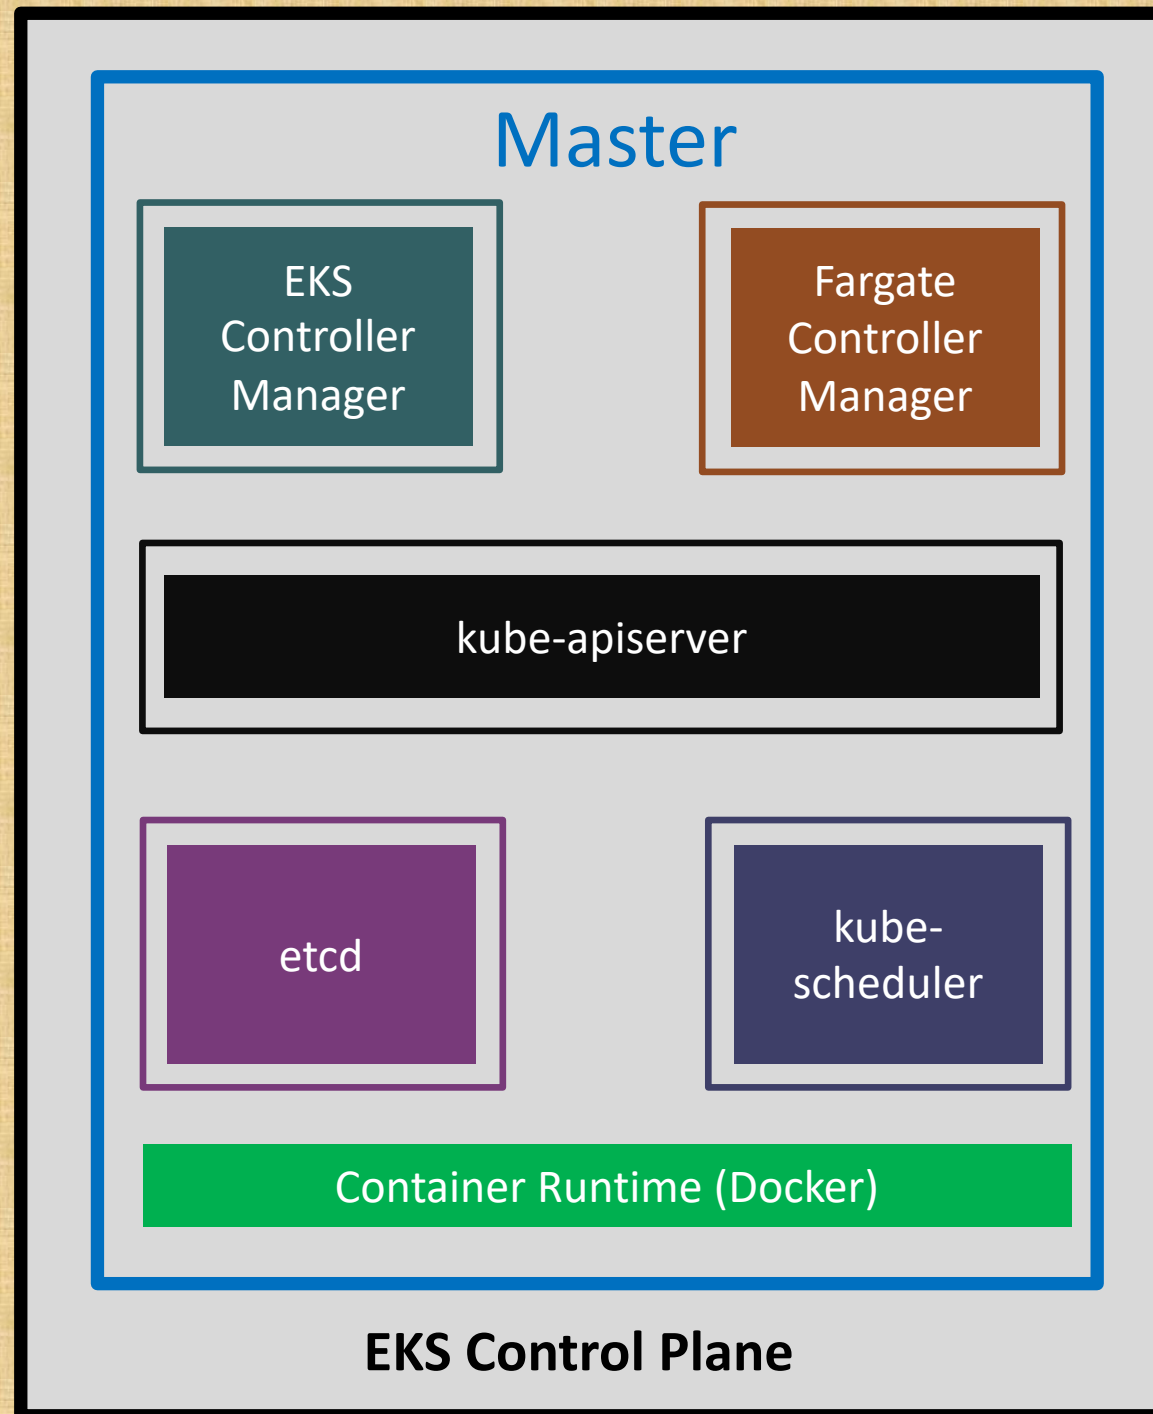
- It is a **network proxy** that runs on each node in your cluster.
- It maintains **network rules** on nodes
- In short, these network rules allow network communication to your Pods from network sessions inside or outside of your cluster.



# AWS EKS Cluster



# EKS Kubernetes - Architecture



# Kubernetes Fundamentals

Pod, ReplicaSet, Deployment & Service



Author: Nho Luong

Skill: DevOps Engineer Lead



# Kubernetes - Fundamentals

k8s Fundamentals

```
graph LR; A[k8s Fundamentals] --> B[Pod]; A --> C[ReplicaSet]; A --> D[Deployment]; A --> E[Service];
```

Pod

A POD is a single instance of an Application.  
A POD is the smallest object, that you can create in Kubernetes.

ReplicaSet

A ReplicaSet will maintain a stable set of replica Pods running at any given time.  
In short, it is often used to guarantee the availability of a specified number of identical Pods

Deployment

A Deployment runs multiple replicas of your application and automatically replaces any instances that fail or become unresponsive. Rollout & rollback changes to applications. Deployments are well-suited for stateless applications.

Service

A service is an abstraction for pods, providing a stable, so called virtual IP (VIP) address.  
In simple terms, service sits Infront of a POD and acts as a load balancer.



# Kubernetes - Imperative & Declarative

## Kubernetes Fundamentals

### Imperative

kubectl

Pod

ReplicaSet

Deployment

Service

### Declarative

YAML & kubectl

Pod

ReplicaSet

Deployment

Service


# Kubernetes

## POD




# Kubernetes - POD

- With Kubernetes our core goal will be to **deploy our applications** in the form of **containers** on **worker nodes** in a k8s cluster.
- Kubernetes **does not** deploy containers directly on the worker nodes.
- Container is **encapsulated** in to a Kubernetes Object named **POD**.
- A POD is a **single instance** of an application.
- A POD is the **smallest object** that we can create in Kubernetes.




Nginx Container  
Image




# Kubernetes - POD

- PODs generally have **one to one** relationship with containers.
- To scale up we **create** new POD and to scale down we **delete** the POD.



# Kubernetes – PODs

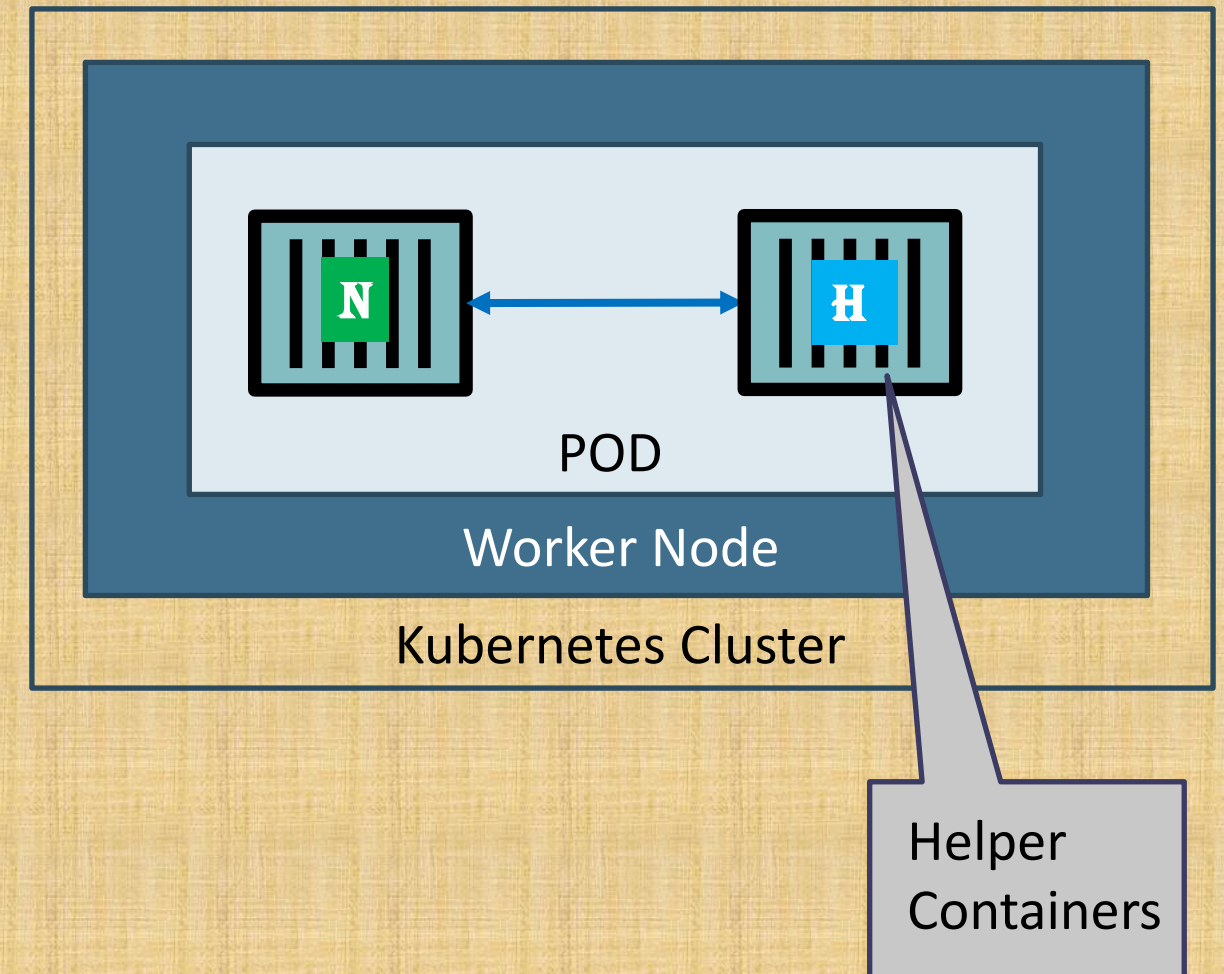
- We **cannot have** multiple containers of **same kind** in a single POD.
- **Example:** Two NGINX containers in single POD serving same purpose is **not recommended**.





# Kubernetes – Multi-Container Pods

- We can have multiple containers in a single POD, provided **they are not of same kind**.
- **Helper Containers (Side-car)**
  - **Data Pullers:** Pull data required by Main Container
  - **Data pushers:** Push data by collecting from main container (logs)
  - **Proxies:** Writes static data to html files using Helper container and Reads using Main Container.
- **Communication**
  - The two containers can easily communicate with each other easily as they share same **network space**.
  - They can also easily share **same storage space**.
- Multi-Container Pods is a **rare use-case** and we will try to focus on core fundamentals.



# Kubernetes

## PODs

## Demo



# Kubernetes Services - NodePort



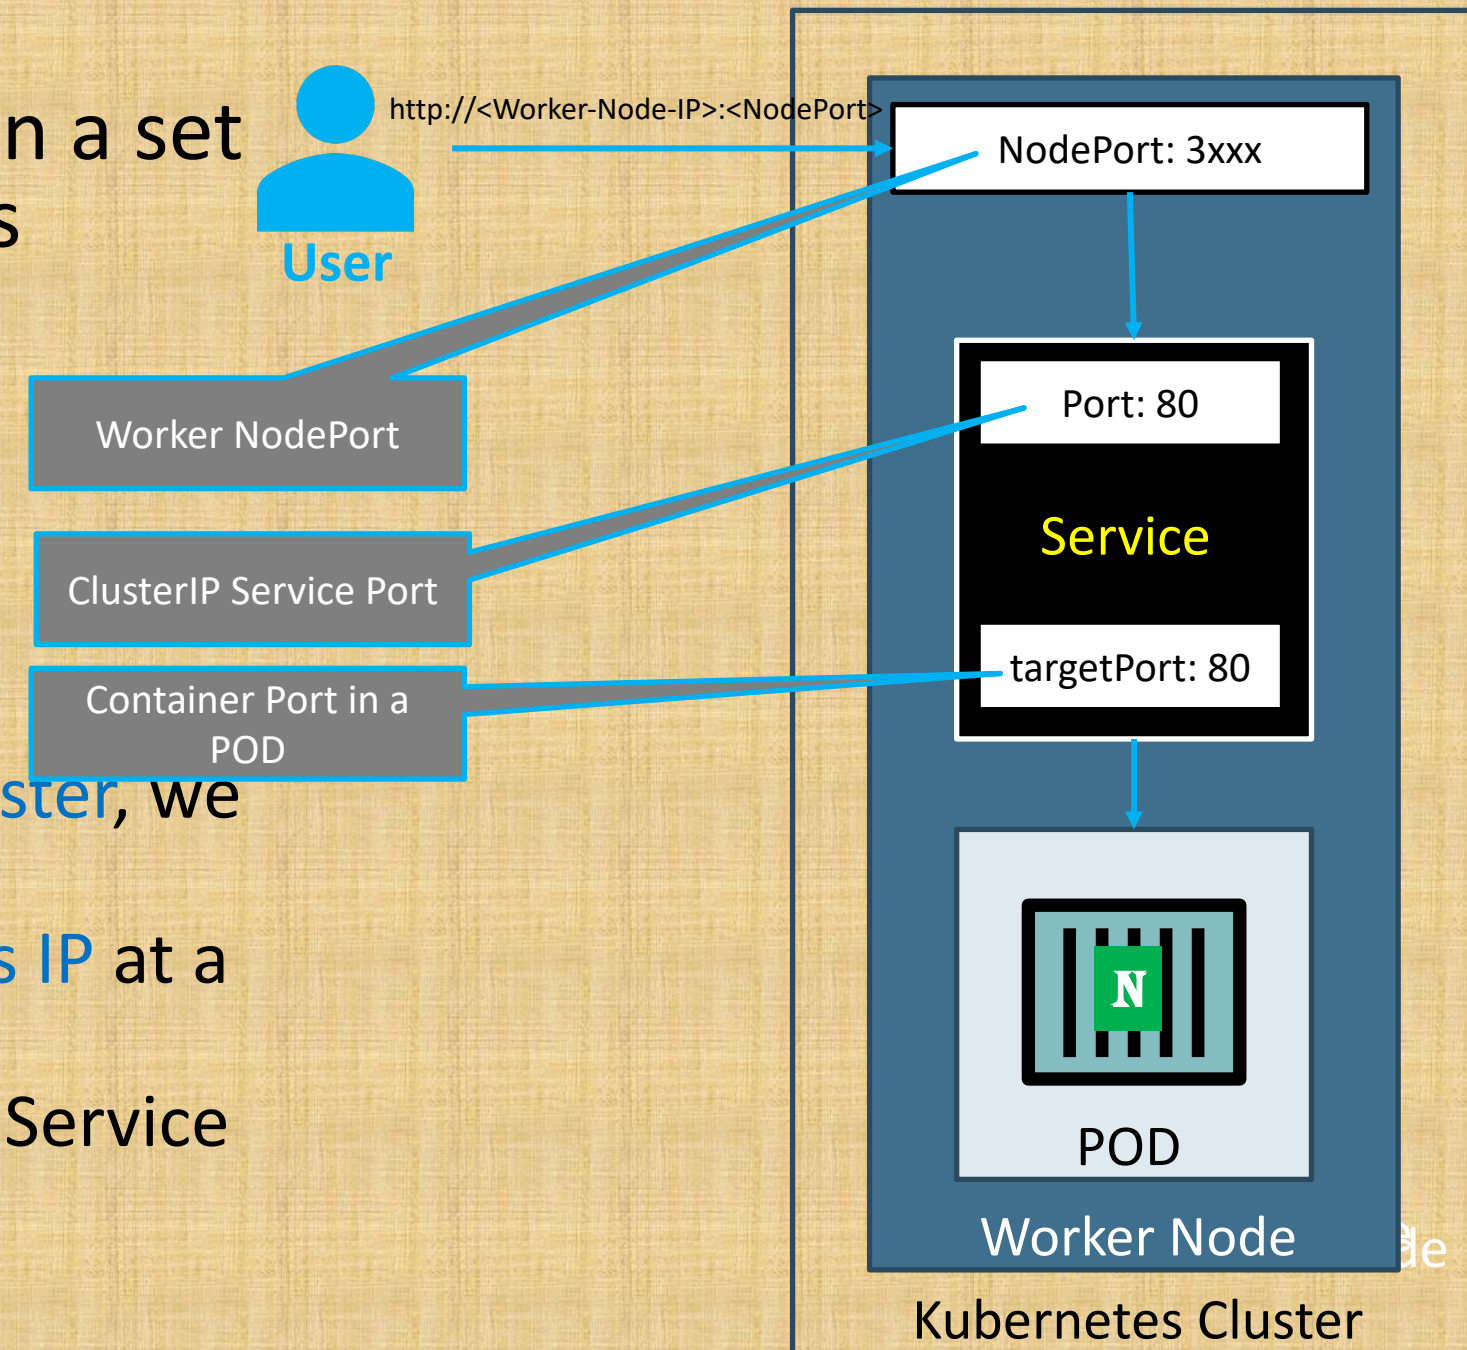
# Kubernetes – Service - NodePort

- We can **expose an application** running on a set of **PODs** using different types of Services available in k8s.

- ClusterIP
- NodePort
- LoadBalancer

- **NodePort Service**

- To access our application **outside of k8s cluster**, we can use NodePort service.
- Exposes the Service on each **Worker Node's IP** at a static port (nothing but NodePort).
- A **ClusterIP** Service, to which the **NodePort** Service routes, is **automatically** created.
- Port Range **30000-32767**





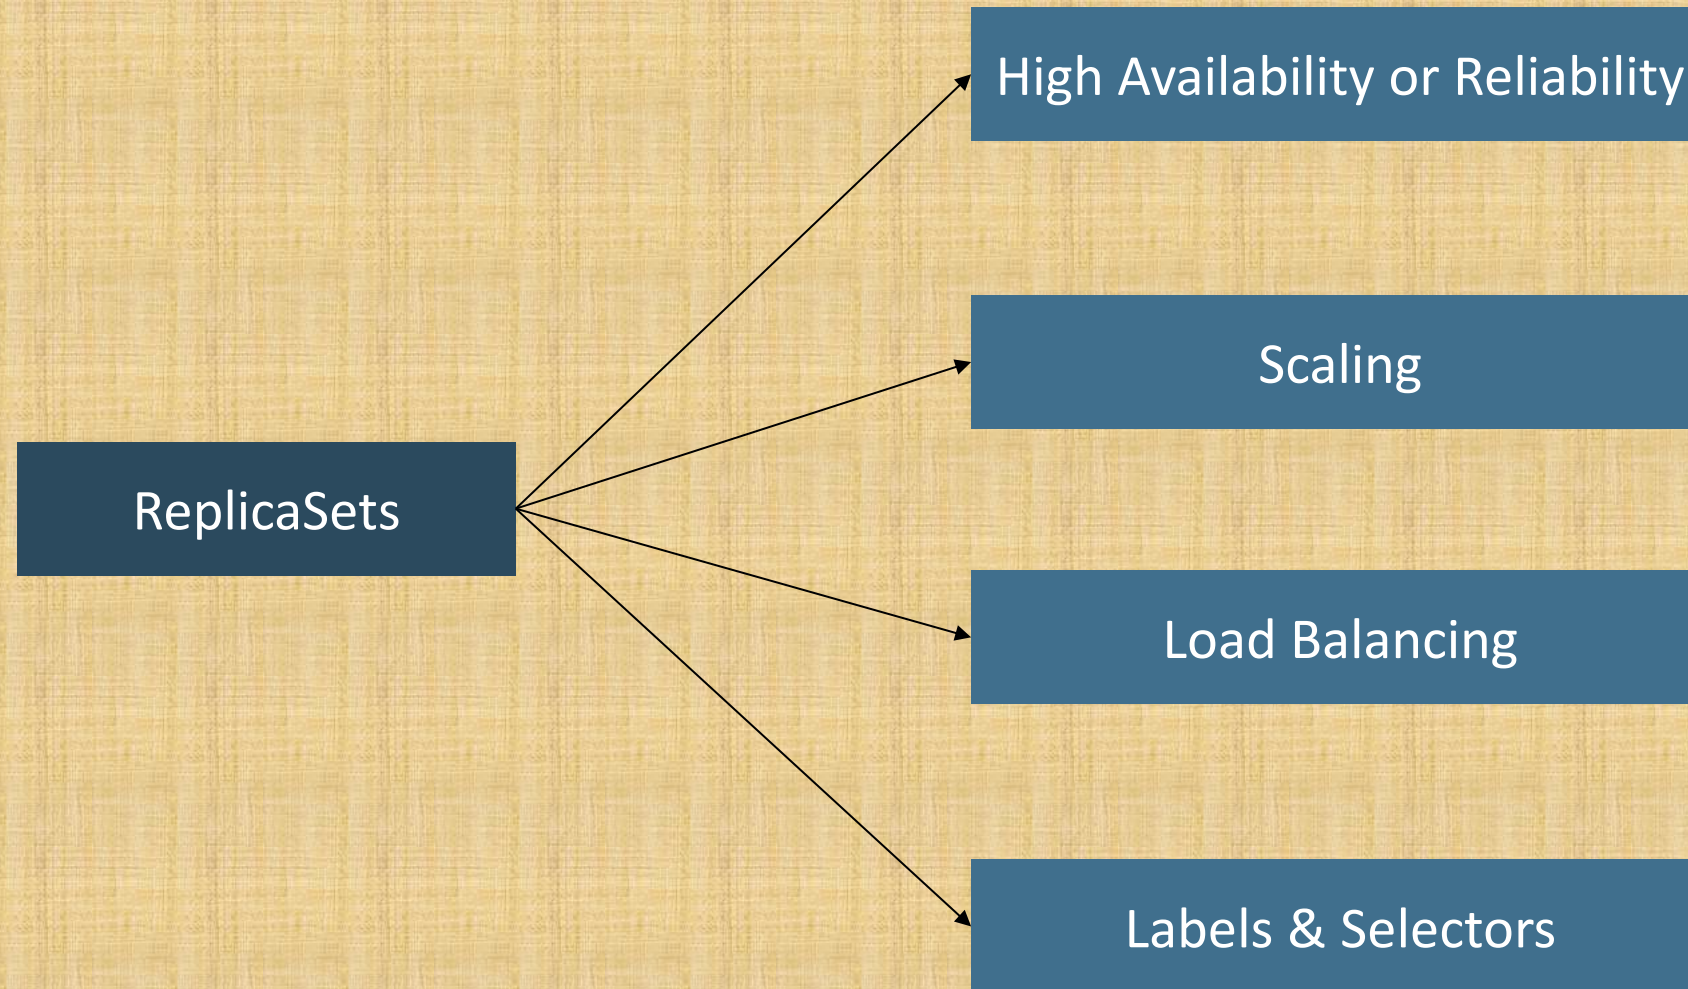
# Kubernetes POD & NodePort Service Demo



# Kubernetes ReplicaSets



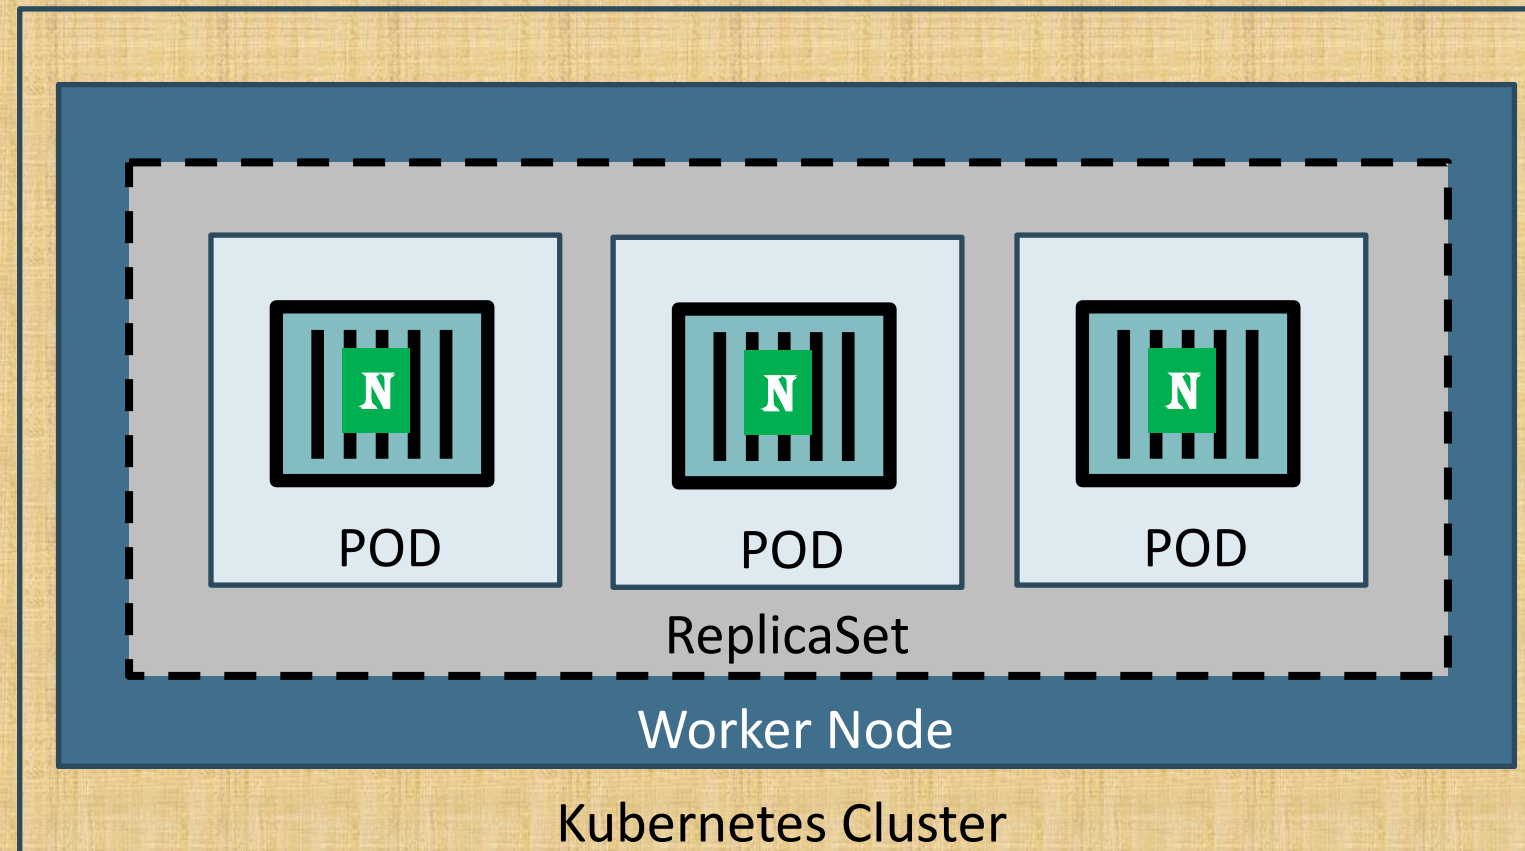
# Kubernetes - ReplicaSets



# Kubernetes – ReplicaSet

- A ReplicaSet's purpose is to maintain a **stable set of replica Pods** running at any given time.
- If our **application crashes (any pod dies)**, replicaset will **recreate** the pod immediately to ensure the configured number of pods running at any given time.


Reliability  
Or  
High Availability






# Kubernetes – ReplicaSet

- Load Balancing
- To avoid overloading of traffic to single pod we can use **load balancing**.
- Kubernetes provides pod load balancing **out of the box** using **Services** for the pods which are part of a ReplicaSet
- **Labels & Selectors** are the **key items** which **ties** all 3 together (Pod, ReplicaSet & Service), we will know in detail when we are writing YAML manifests for these objects



# Kubernetes – ReplicaSet

- Scaling
- When load become too much for the number of existing pods, Kubernetes enables us to easily **scale** up our application, adding additional pods as needed.
- This is going to be **seamless and super quick**.



# Kubernetes ReplicaSets Demo

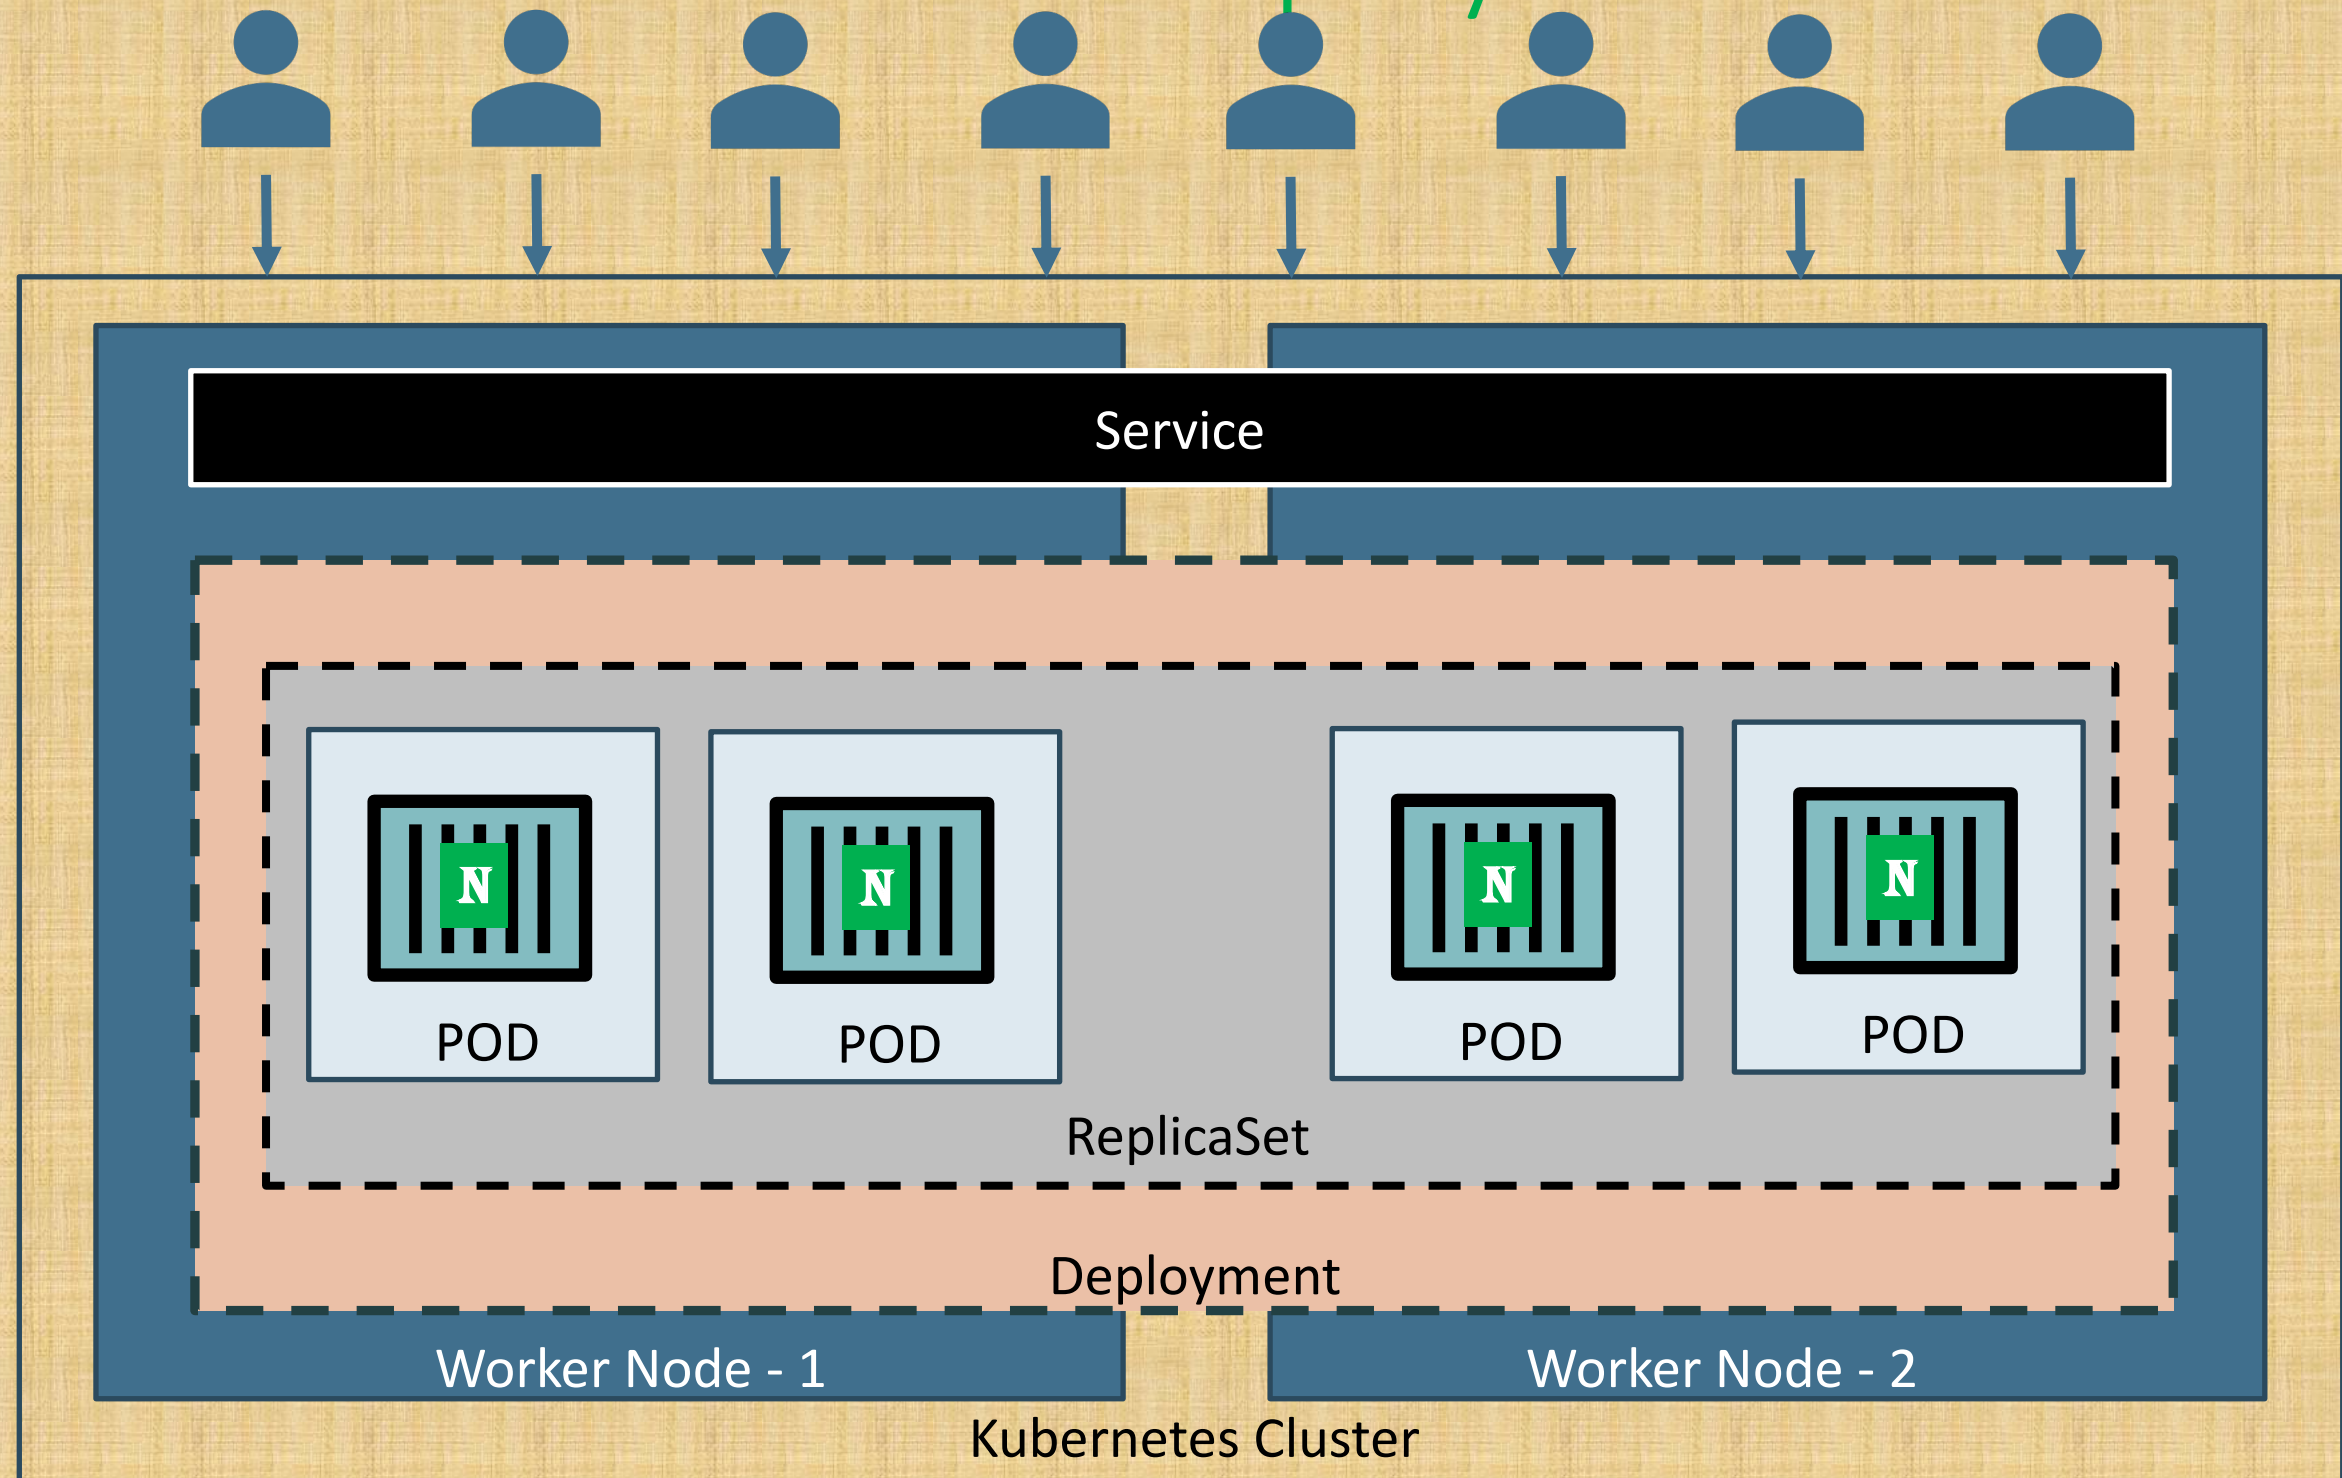


# Kubernetes Deployments

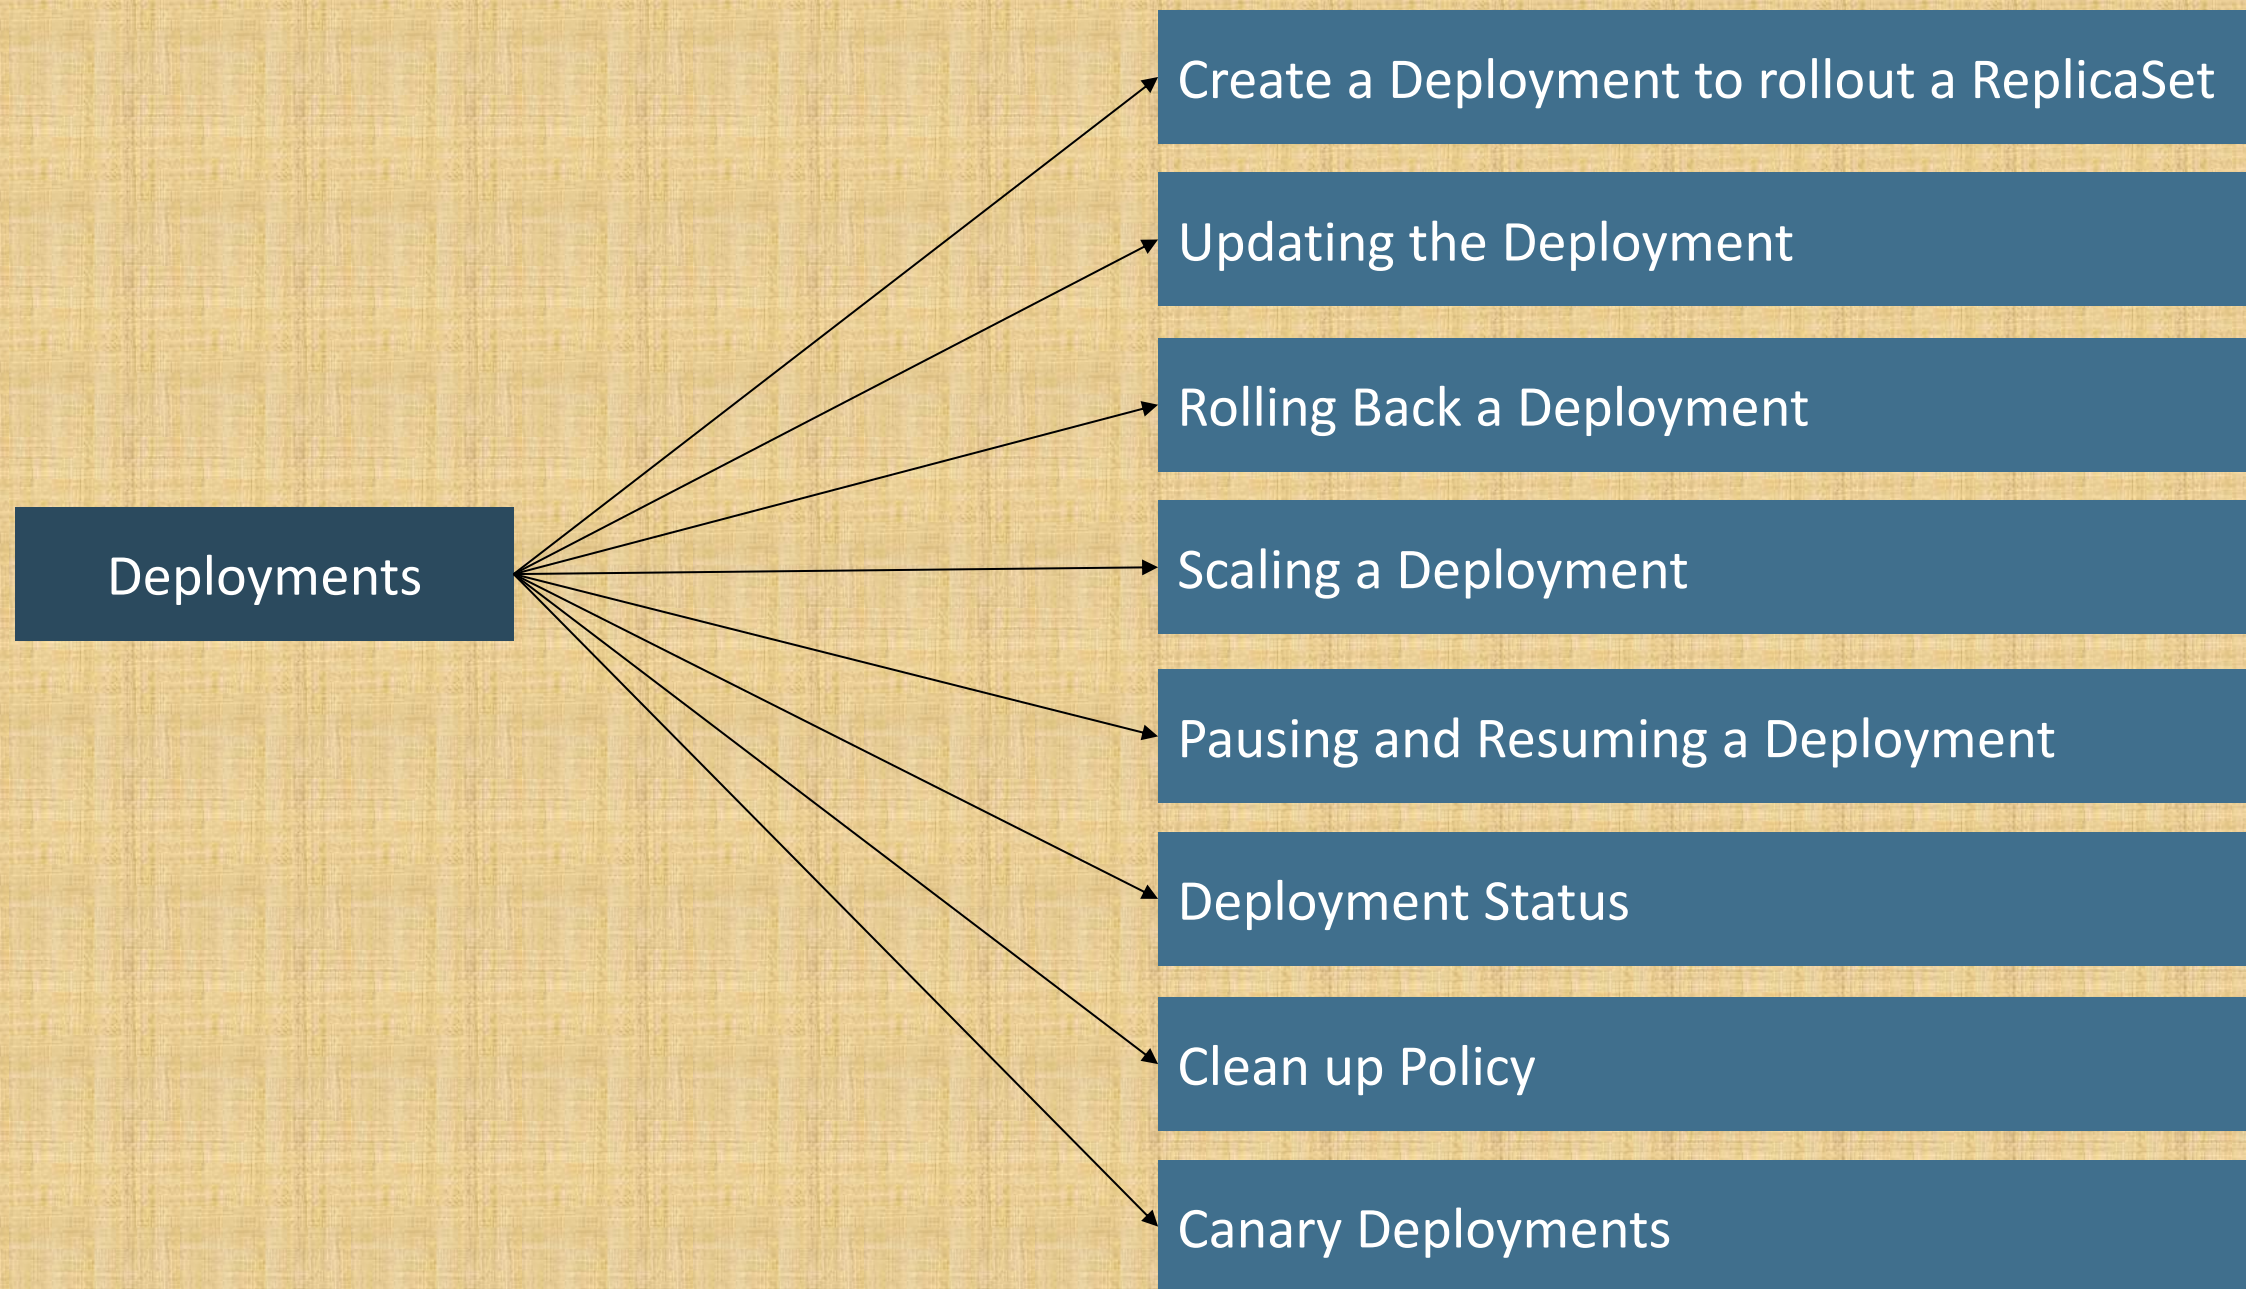




# Kubernetes – Deployments



# Kubernetes - Deployment



# Kubernetes Deployments Demo

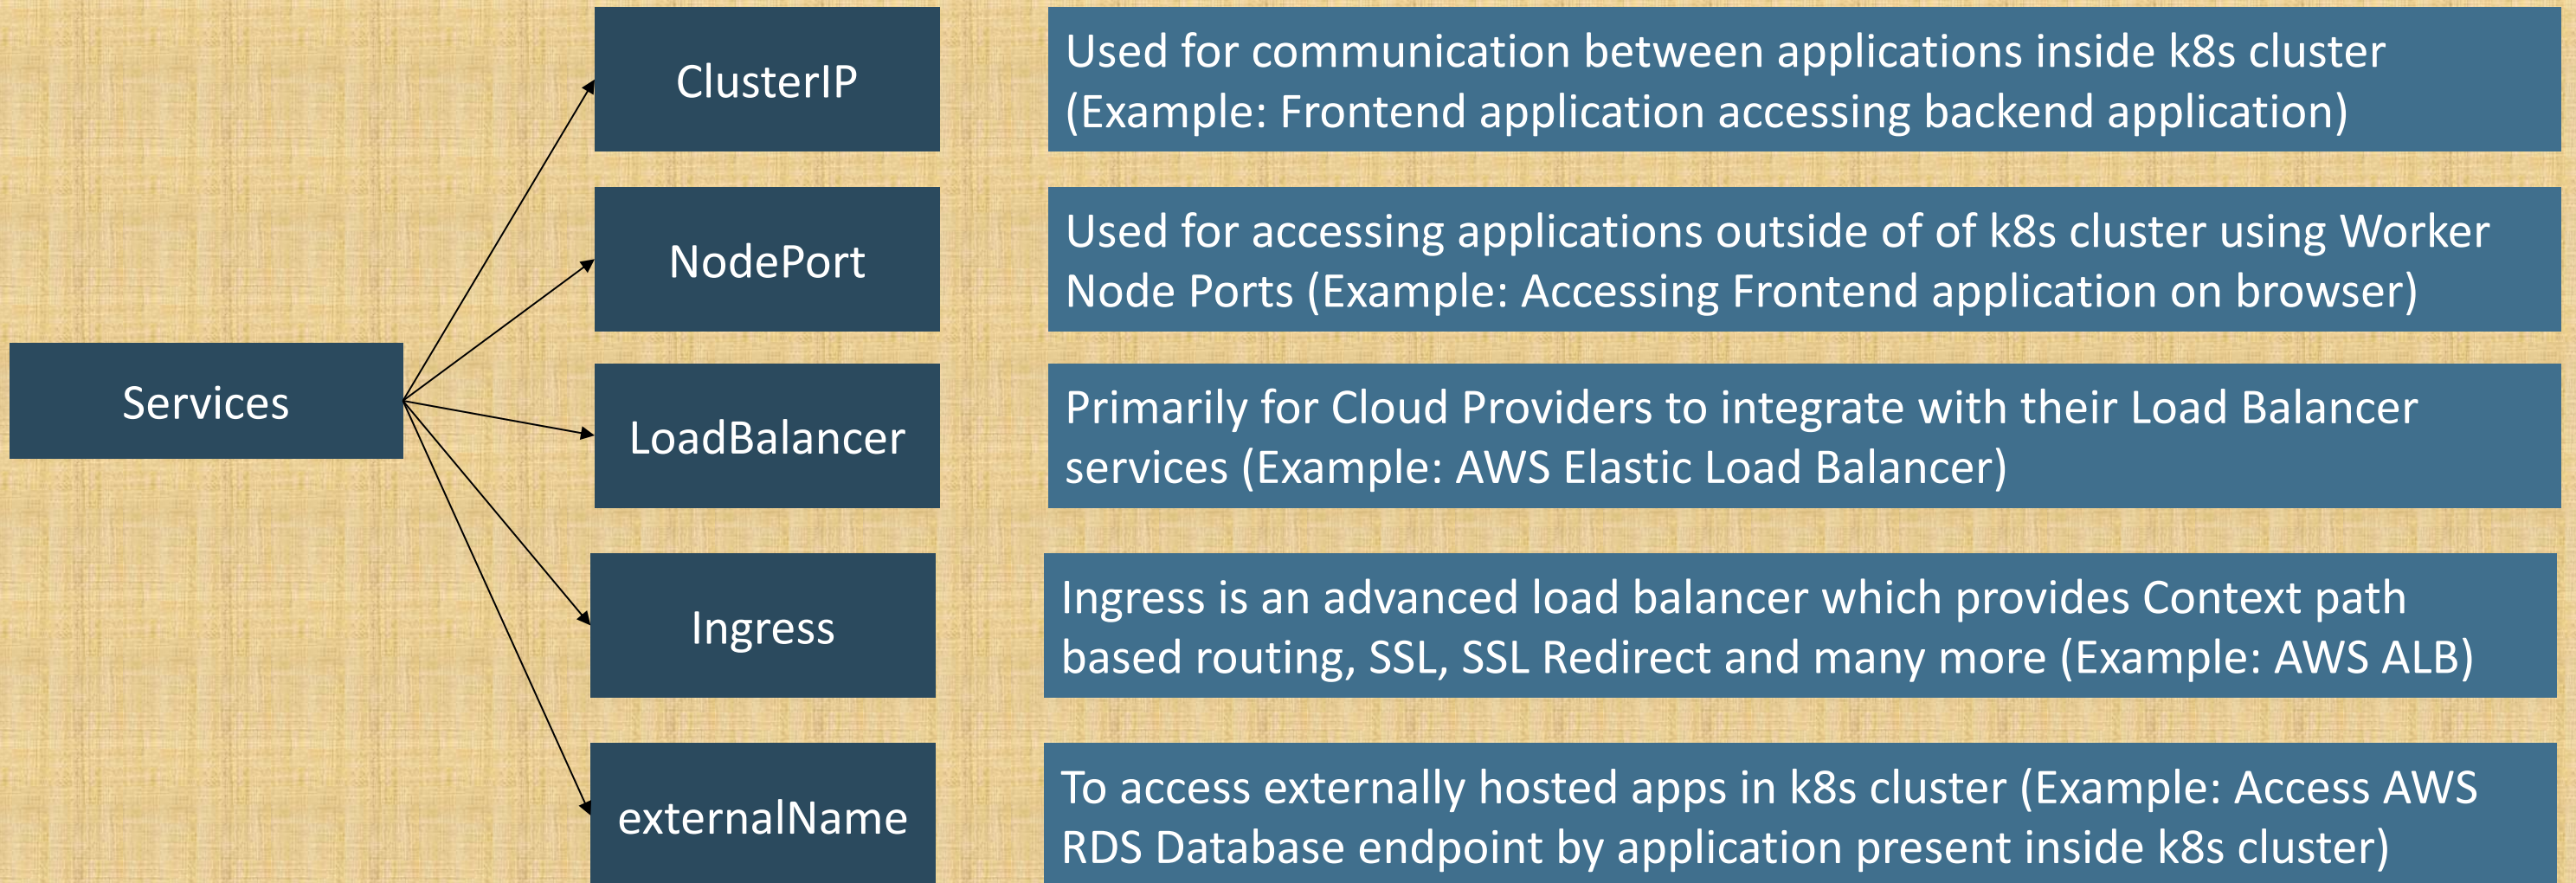



# Kubernetes Services






# Kubernetes - Services






Users


Frontend App – NodePort or LoadBalancer or Ingress Service



Backend App - ClusterIP Service



DB – ExternalName Service



# Services

Author: Nho Luong

Skill: DevOps Engineer Lead

# Kubernetes Services Demo

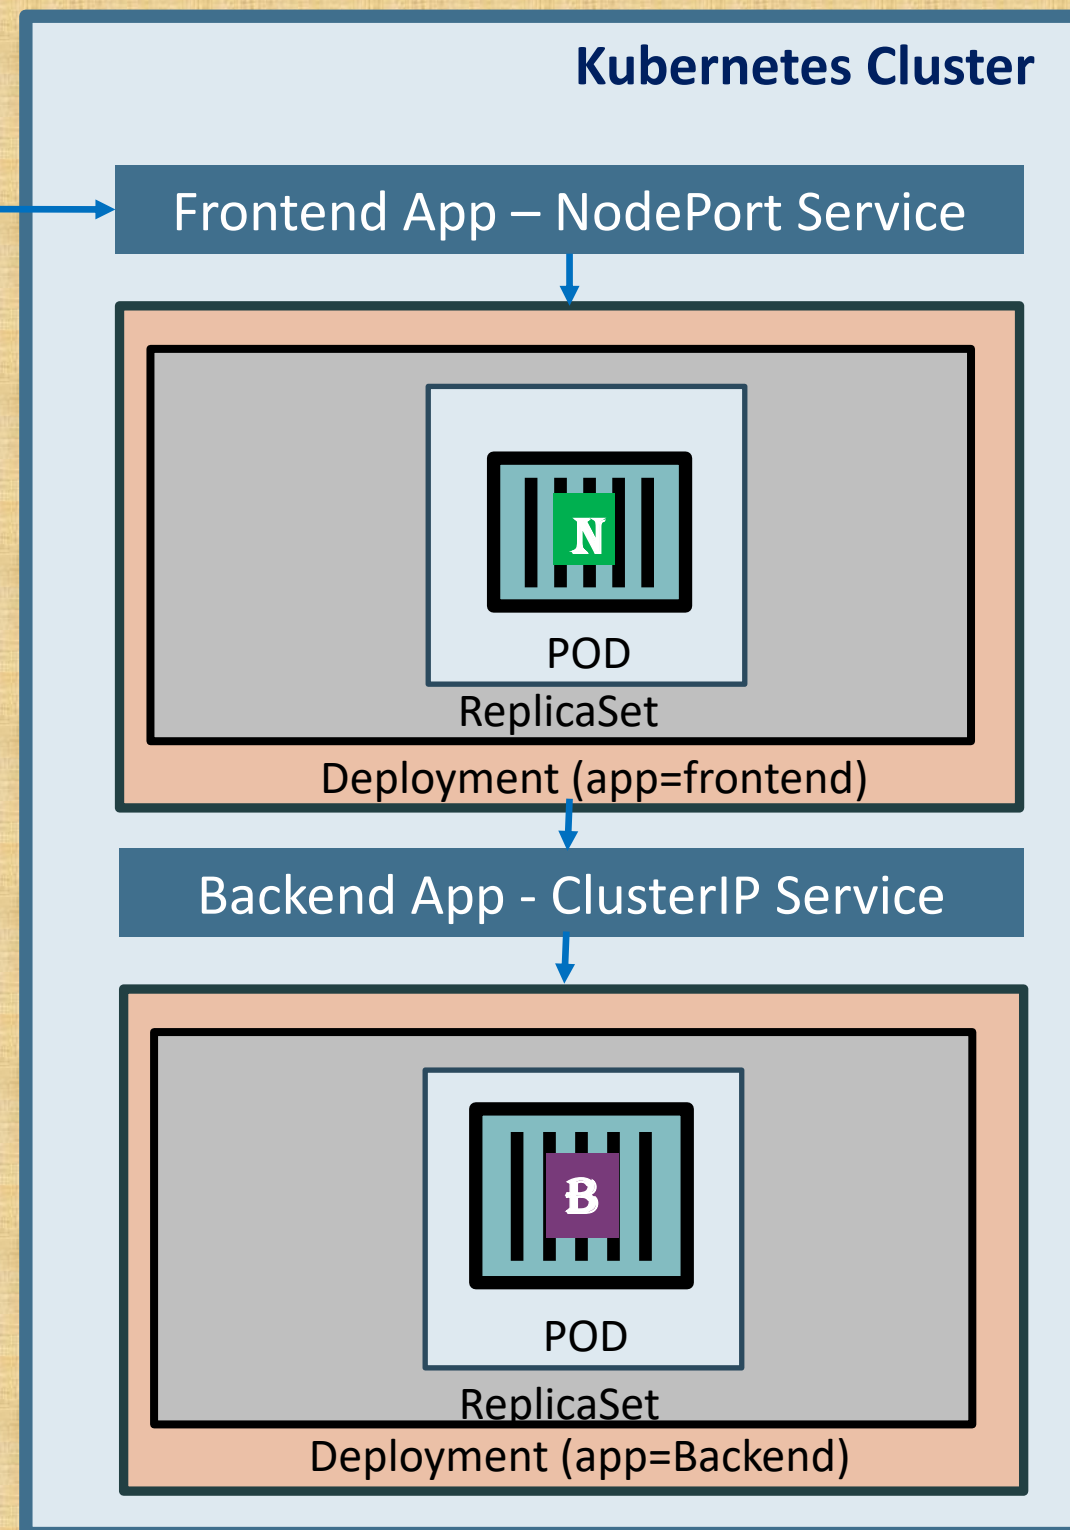


# Services Demo



Users

`http://<workernode-public-ip>:<NodePort>/hello`





# Kubernetes

## YAML Basics



# YAML Basics

- YAML is **not a** Markup Language
- YAML is used to **store information** about different things
- We can use YAML to **define key, Value pairs** like variables, lists and objects
- YAML is very similar to **JSON** (Javascript Object Notation)
- YAML primarily focuses on **readability** and **user friendliness**
- YAML is designed to be **clean and easy to read**
- We can define YAML files with two different extensions
  - abc.**yml**
  - abc.**yaml**

# YAML Basics

- YAML Comments
- YAML Key Value Pairs
- YAML Dictionary or Map
- YAML Array / Lists
- YAML Spaces
- YAML Document Separator



Thank You