

Kubernetes Orchestration

Author: Nho Luong
Skill: DevOps Engineer Lead

Microsoft

Nho Luong

has successfully passed all requirements for

Microsoft Certified: Azure Solutions Architect Expert

Credential ID: 3CF3D14D253FE551
Certification number: 46F8FB-7D3F5B
Earned on: June 14, 2024
Expires on: June 15, 2025

Microsoft Certified

EXPERT


Satya Narayana Nadella

Online Verifiable

aws certified

nho luong

AWS Certified Solutions Architect - Professional

VALIDATION NUMBER: 119ee670510e4bb99fe2956152effb7e
VALIDATE AT: <https://aws.amazon.com/verification>

Issue Date: June 21, 2024
Expiration Date: June 21, 2027



Alpine

aws

HOME

PROFILE

EXAM REGISTRATION

EXAM HISTORY

CERTIFICATIONS

Certification status

Agreement

BENEFITS

DIGITAL BADGES

SUPPORT AND FAQs

Certification status

NHO LUONG

AWS04440050

Active

PROFESSIONAL

AWS Certified Solutions Architect - Professional

SAP

ACTIVE DATE
2024-06-21

EXPIRATION DATE
2027-06-21

VIEW MORE

Active

SPECIALTY

AWS Certified Security - Specialty

SCS

ACTIVE DATE
2024-06-19

EXPIRATION DATE
2027-06-19

VIEW MORE

Active

PROFESSIONAL

AWS Certified DevOps Engineer - Professional

DOP

ACTIVE DATE
2024-06-17

EXPIRATION DATE
2027-06-17

VIEW MORE

Active

ASSOCIATE

AWS Certified Solutions Architect - Associate

SAA

ACTIVE DATE
2024-06-15

EXPIRATION DATE
2027-06-21

VIEW MORE

Active

FOUNDATIONAL

AWS Certified Cloud Practitioner

CLF

ACTIVE DATE
2024-05-27

EXPIRATION DATE
2027-06-21

Activity

Training

Plans

Challenges

Credentials

Q&A

Achievements

Collections

Transcript

4 items

Sorted by expiration date

CERTIFICATION

Microsoft Certified: DevOps Engineer Expert

Expires on June 15, 2025 at 6:59 AM (UTC +07:00) • Earned on June 14, 2024

View certification details

CERTIFICATION

Microsoft Certified: Azure Solutions Architect Expert

Expires on June 15, 2025 at 6:59 AM (UTC +07:00) • Earned on June 14, 2024

View certification details

CERTIFICATION

Microsoft Certified: Azure Administrator Associate

Expires on June 15, 2025 at 6:59 AM (UTC +07:00) • Earned on June 14, 2024

View certification details


CERTIFICATION

Microsoft Certified: Azure Fundamentals

Earned on May 24, 2024

View certification details

Securing Kubernetes Workloads



Best Practices for Securing Kubernetes Workload Configurations Across Clouds

Kubernetes Security Framework

	Build	Operate
Container Hosts:	<ul style="list-style-type: none">Minimal OSOS Hardening	<ul style="list-style-type: none">CIS Benchmarks
Clusters:	<ul style="list-style-type: none">RBACAudit Policies and LoggingCertificate Management	<ul style="list-style-type: none">Identity and AccessKubernetes upgradesCIS Benchmarks
Applications:	<ul style="list-style-type: none">Image scanning	<ul style="list-style-type: none">Image ProvenanceSecrets ManagementNamespacesAccess ControlsNetwork PoliciesResource QuotasPod Security Policy

Kubernetes Security Framework

	Build	Operate
Container Hosts:	<ul style="list-style-type: none">Minimal OSOS Hardening	<ul style="list-style-type: none">CIS Benchmarks
Clusters:	<ul style="list-style-type: none">RBACAudit Policies and LoggingCertificate Management	<ul style="list-style-type: none">Identity and AccessKubernetes upgradesCIS Benchmarks
Applications:	<ul style="list-style-type: none">Image scanning	<ul style="list-style-type: none">Image ProvenanceSecrets ManagementNamespacesAccess ControlsNetwork PoliciesResource QuotasPod Security Policy

Image Provenance

- Image scanning checks images for vulnerabilities
 - o Ideally done when the image is built and before it is accepted into the image registry
- Image provenance
 1. Confirms that an image being deployed is from a trusted source
 2. Confirms that image has not been not tampered with

Image Provenance - Solutions

- **Kubernetes ImagePolicyWebhook**
 - Configured as an admission controller
 - Sends an ImageReview request
 - Expects an ImageReview response of accept or deny


Image Provenance - Solutions

- **Portieris**

- Also an admission controller
- Integrates with Notary (a content trust store) – part of the The Update Framework (TUF)
- Provides way to specify image security policies at a namespace and cluster level

Image Provenance – Partial Solutions

- **Kyverno**
 - Also an admission controller
 - Kubernetes Native Policy Engine
 - Policies are written as overlay rules



```
kind: ClusterPolicy
metadata:
  name: validate-image-registry
spec:
  rules:
  - name: validate-image-registry
    match:
      resources:
        kinds:
        - Pod
    validate:
      message: "Image registry is not allowed"
      pattern:
        spec:
          containers:
          - name: "*"
            image: !image
```


Image Provenance – Partial Solutions

- **OPA / Gatekeeper**
 - Also an admission controller
 - General Purpose Policy Engine
 - Policies are written in Rego

```
package kubernetes.admission

import data.kubernetes.namespaces

deny[msg] {
    input.request.kind.kind = "Deployment"
    input.request.operation = "CREATE"
    registry =
input.request.object.spec.template.spec.containers[_].image
    name = input.request.object.metadata.name
    namespace = input.request.object.metadata.namespace
    not reg_matches_any(registry, valid_deployment_registries)
    msg = sprintf("invalid deployment, namespace=%q, name=%q, registry=%q", [namespace, name, registry])
}

valid_deployment_registries = {registry |
    whitelist = "<COMMA_SEPARATED_LIST_OF_ALLOWED_REGISTRIES>"
    registries = split(whitelist, ",")
    registry = registries[_]
}

reg_matches_any(str, patterns) {
    reg_matches(str, patterns[_])
}
```

Secrets Management Anti-Patterns

(please try not to do this)

x Hard-coded

x Packaged with code

x Inserted via build tools x

Environment Variables

- Any sensitive data that an application needs
 - Passwords
 - Certificates
 - Keys
 - ...

What Kubernetes Provides

- API Object to define secrets
- Values are base 64 encoded (default)
- Secrets are namespaced
- Secrets can be mounted as volumes
- Secrets can be used as environment variables
- Encryption can be configured at the API Server

```
apiVersion: v1 kind:  
Secret metadata: name:  
mysecret type: Opaque  
data:  
username: YWRtaW4=
```


So, what's missing?

Kubernetes secrets are a step forward, but have a few limitations:

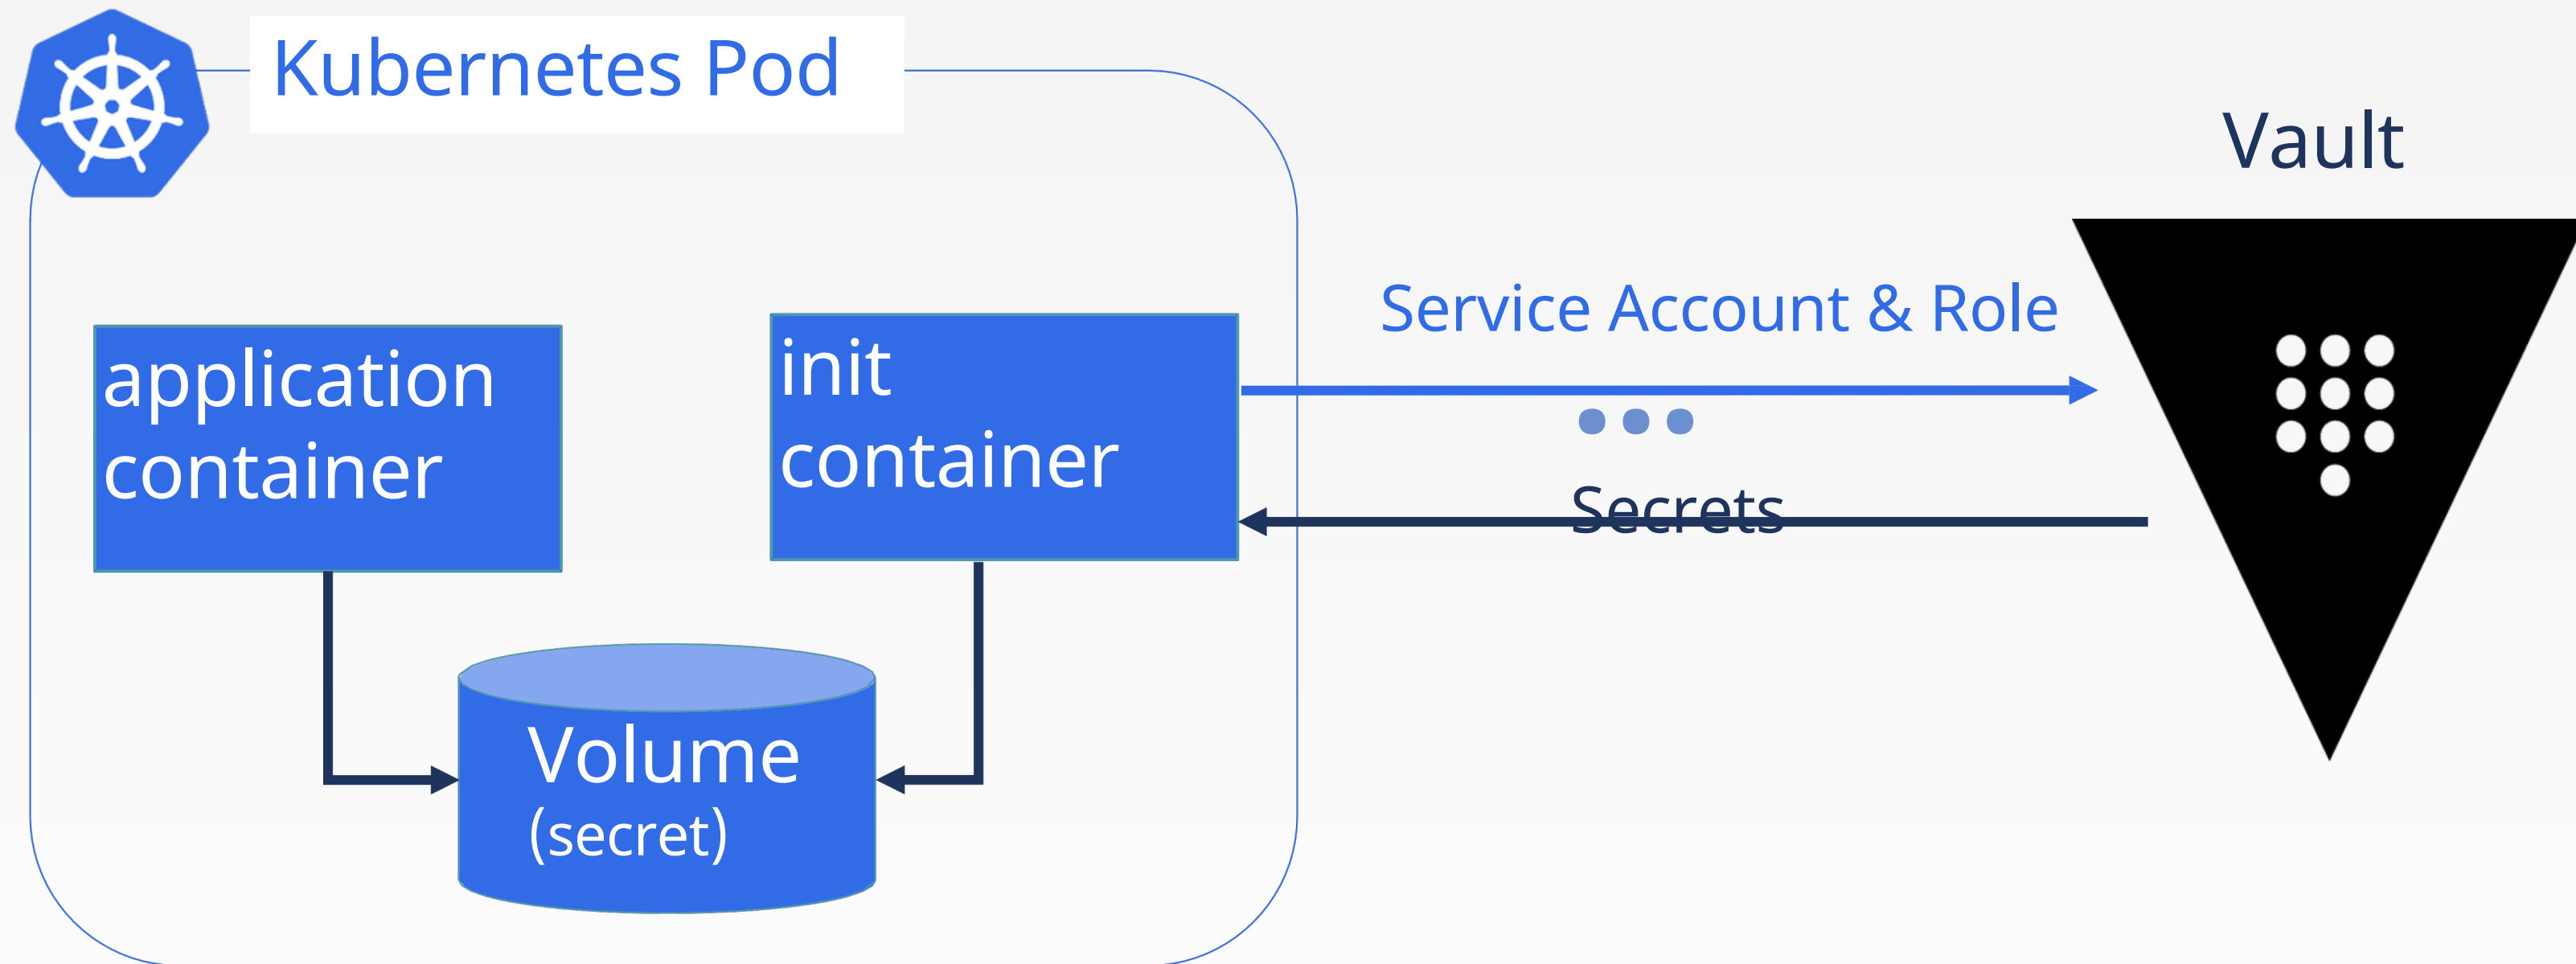
- Encryption requires configuring static keys or a KMS
- Shared (static) approach
- No leases, rotation, etc.

Secrets Management with Hashicorp Vault

- Helps automates security best practices for
 - Secrets Management
 - Auditing
 - Certificate Management
 - Encryption
- Dynamic Secrets
 - Credentials (keys, passwords, certificates) are generated when a client requests them
 - Credentials are per client
 - Credentials are automatically deleted if a lease expires



An init container to fetch secrets



Namespaces

- **Kubernetes Data Plane Virtualization**

Kubernetes supports multiple virtual clusters backed by the same physical cluster. These virtual clusters are called namespaces.

- Namespaces partition the Kubernetes object model so multiple objects with the same name can exist in the same cluster
- Namespaces are the foundation for applying other security constructs

Role-based access control (RBAC)

- Users are authenticated via OIDC, X.509 certificates, tokens, etc.
- The authentication result can provide user and group information.
- However, Users and User Groups are managed externally (e.g. in an LDAP / AD server).
- Kubernetes has a fine-grained permission model
 - Role (namespace) / ClusterRole
- Roles are mapped to users or groups via role bindings
 - RoleBinding (namespace) / ClusterRoleBinding

Service Accounts

- Service Accounts are meant for authenticating and authorizing processes
- Each namespace has a default service account
- Each Pod has a service account (default – if not specified)
- A best practice is to use a service account per app
- To prevent a service account token from being mounted in a Pod use “**automountServiceAccountToken: false**”. This can be enforced via a policy.

Network Segmentation via Network Policies

- By default, Kubernetes pods are **“non-isolated”**
 - They accept network connections from any source and can initiate connection requests to any destination
- Network Policies define traffic rules for Kubernetes pods
 - ingress (inbound traffic)
 - egress (outbound traffic)

Network Policy

Pod Selector

Ingress



Ingress Rule

Egress



Egress Rule

Resource Management


- Pods can have resource requests and limits
- This allows three quality of service models



Burstable



Guaranteed



- A namespace can have limits and default allocations
- Quotas and limits ensure fairness and stability

Pod Security Policies

- Controls runtime security settings for pods
- Enabled at the API Controller
- Requires a role binding between pod Service Account and the PSP

Control Aspect	Field Names
Running of privileged containers	privileged
Usage of host namespaces	hostPID , hostIPC
Usage of host networking and ports	hostNetwork , hostPorts
Usage of volume types	volumes
Usage of the host filesystem	allowedHostPaths
White list of Flexvolume drivers	allowedFlexVolumes
Allocating an FSGroup that owns the pod's volumes	fsGroup
Requiring the use of a read only root file system	readOnlyRootFilesystem
The user and group IDs of the container	runAsUser , runAsGroup , supplementalGroups
Restricting escalation to root privileges	allowPrivilegeEscalation , defaultAllowPrivilegeEscalation
Linux capabilities	defaultAddCapabilities , requiredDropCapabilities , allowedCapabilities
The SELinux context of the container	seLinux
The Allowed Proc Mount types for the container	allowedProcMountTypes
The AppArmor profile used by containers	annotations
The seccomp profile used by containers	annotations
The sysctl profile used by containers	forbiddenSysctls , allowedUnsafeSysctls

Use a policy engine to audit and enforce

- Pod Security Policies are tricky to manage
 - Require a role binding to SA
 - Applied in alphabetical order
- Kyverno supports enforcement of the important PSP checks

```
kind: ClusterPolicy
metadata:
  name: validate-deny-runasrootuser
spec:
  validationFailureAction: "audit"
  rules:
  - name: deny-runasrootuser
    exclude:
      resources:
        namespaces:
        - kube-system
    match:
      resources:
        kinds:
        - Pod
    validate:
      message: "Root user is not allowed. Set runAsNonRoot to true."
      anyPattern:
      - spec:
          securityContext:
            runAsNonRoot: true
      - spec:
          containers:
          - name: "*"
            securityContext:
```



Thank You