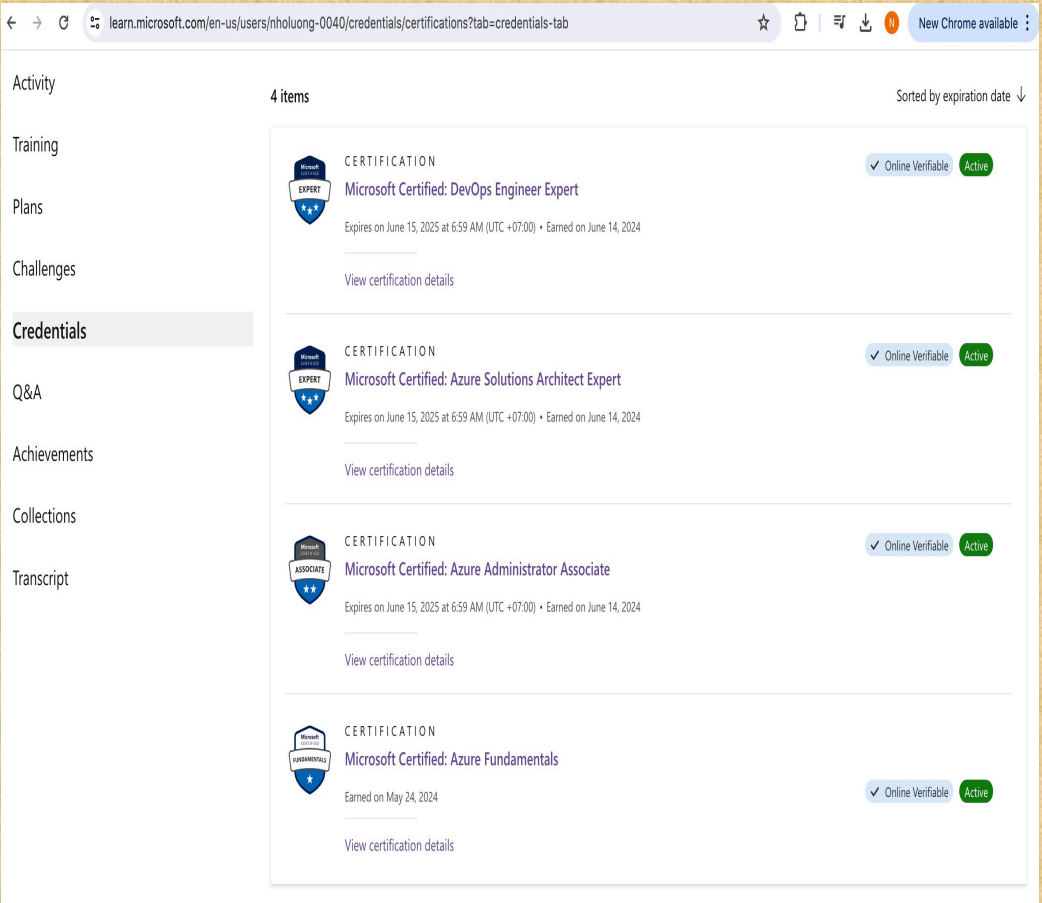
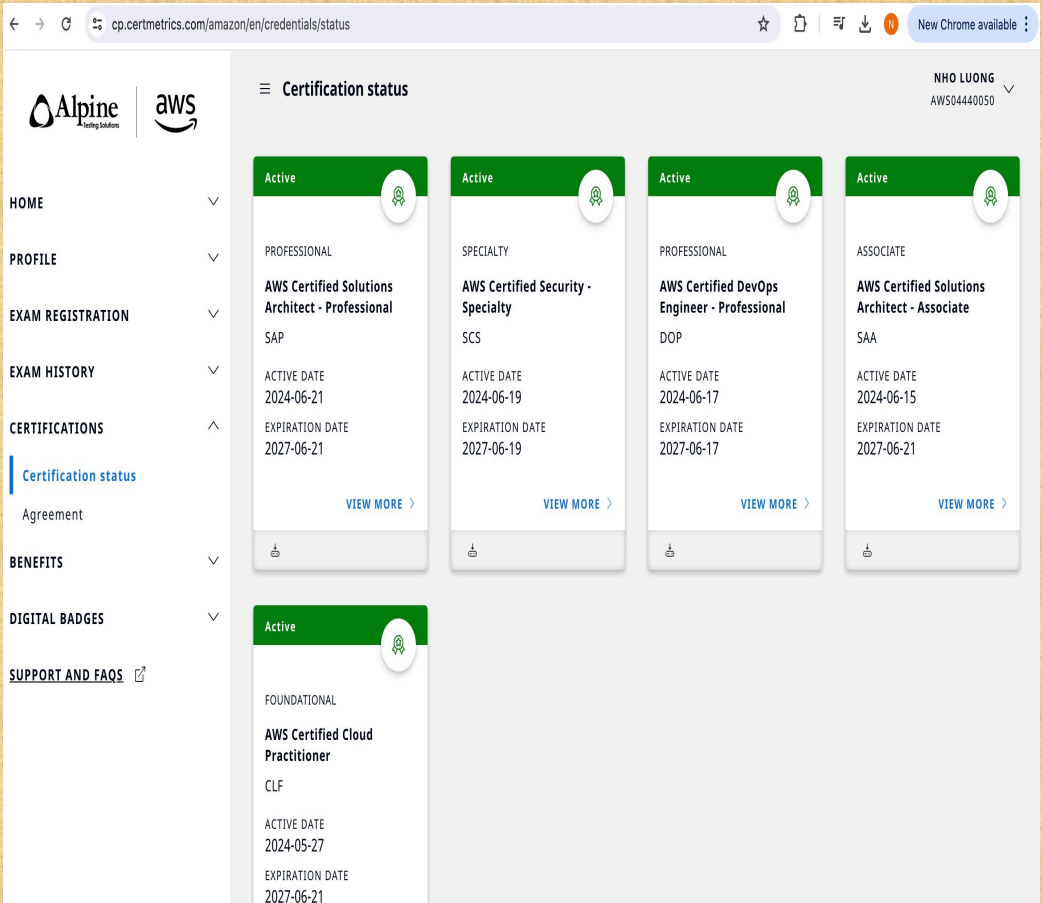
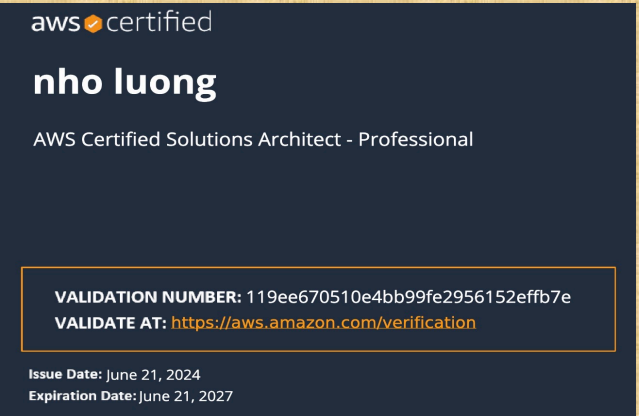



# Kubernetes Orchestration

Author: Nho Luong  
Skill: DevOps Engineer Lead



# Securing Kubernetes Workloads



*Best Practices for Securing Kubernetes Workload Configurations Across Clouds*

# Kubernetes Security Framework

	Build	Operate
Container Hosts:	<ul style="list-style-type: none"><li>Minimal OS</li><li>OS Hardening</li></ul>	<ul style="list-style-type: none"><li>CIS Benchmarks</li></ul>
Clusters:	<ul style="list-style-type: none"><li>RBAC</li><li>Audit Policies and Logging</li><li>Certificate Management</li></ul>	<ul style="list-style-type: none"><li>Identity and Access</li><li>Kubernetes upgrades</li><li>CIS Benchmarks</li></ul>
Applications:	<ul style="list-style-type: none"><li>Image scanning</li></ul>	<ul style="list-style-type: none"><li>Image Provenance</li><li>Secrets Management</li><li>Namespaces</li><li>Access Controls</li><li>Network Policies</li><li>Resource Quotas</li><li>Pod Security Policy</li></ul>

# Kubernetes Security Framework

	Build	Operate
Container Hosts:	<ul style="list-style-type: none"><li>Minimal OS</li><li>OS Hardening</li></ul>	<ul style="list-style-type: none"><li>CIS Benchmarks</li></ul>
Clusters:	<ul style="list-style-type: none"><li>RBAC</li><li>Audit Policies and Logging</li><li>Certificate Management</li></ul>	<ul style="list-style-type: none"><li>Identity and Access</li><li>Kubernetes upgrades</li><li>CIS Benchmarks</li></ul>
Applications:	<ul style="list-style-type: none"><li>Image scanning</li></ul>	<ul style="list-style-type: none"><li>Image Provenance</li><li>Secrets Management</li><li>Namespaces</li><li>Access Controls</li><li>Network Policies</li><li>Resource Quotas</li><li>Pod Security Policy</li></ul>



# Image Provenance

- Image scanning checks images for vulnerabilities
  - o Ideally done when the image is built and before it is accepted into the image registry
- Image provenance
  1. Confirms that an image being deployed is from a trusted source
  2. Confirms that image has not been not tampered with

# Image Provenance - Solutions

- **Kubernetes ImagePolicyWebhook**
  - Configured as an admission controller
  - Sends an ImageReview request
  - Expects an ImageReview response of accept or deny


# Image Provenance - Solutions

- **Portieris**

- Also an admission controller
- Integrates with Notary (a content trust store) – part of the The Update Framework (TUF)
- Provides way to specify image security policies at a namespace and cluster level

# Image Provenance – Partial Solutions

- **Kyverno**
  - Also an admission controller
  - Kubernetes Native Policy Engine
  - Policies are written as overlay rules



```
kind: ClusterPolicy
metadata:
  name: validate-image-registry
spec:
  rules:
  - name: validate-image-registry
    match:
      resources:
        kinds:
        - Pod
    validate:
      message: "Image registry is not allowed"
      pattern:
        spec:
          containers:
          - name: "*"
            image: !image
```



# Image Provenance – Partial Solutions

- **OPA / Gatekeeper**
  - Also an admission controller
  - General Purpose Policy Engine
  - Policies are written in Rego

```
package kubernetes.admission

import data.kubernetes.namespaces

deny[msg] {
    input.request.kind.kind = "Deployment"
    input.request.operation = "CREATE"
    registry =
input.request.object.spec.template.spec.containers[_].image
    name = input.request.object.metadata.name
    namespace = input.request.object.metadata.namespace
    not reg_matches_any(registry, valid_deployment_registries)
    msg = sprintf("invalid deployment, namespace=%q, name=%q, registry=%q", [namespace, name, registry])
}

valid_deployment_registries = {registry |
    whitelist = "<COMMA_SEPARATED_LIST_OF_ALLOWED_REGISTRIES>"
    registries = split(whitelist, ",")
    registry = registries[_]
}

reg_matches_any(str, patterns) {
    reg_matches(str, patterns[_])
}
```

# Secrets Management Anti-Patterns

*(please try not to do this)*

x Hard-coded

x Packaged with code

x Inserted via build tools x

Environment Variables

- Any sensitive data that an application needs
  - Passwords
  - Certificates
  - Keys
  - ...

# What Kubernetes Provides

- API Object to define secrets
- Values are base 64 encoded (default)
- Secrets are namespaced
- Secrets can be mounted as volumes
- Secrets can be used as environment variables
- Encryption can be configured at the API Server

```
apiVersion: v1 kind:  
Secret metadata: name:  
mysecret type: Opaque  
data:  
username: YWRtaW4=
```

# So, what's missing?


*Kubernetes secrets are a step forward, but have a few limitations:*

- Encryption requires configuring static keys or a KMS
- Shared (static) approach
- No leases, rotation, etc.

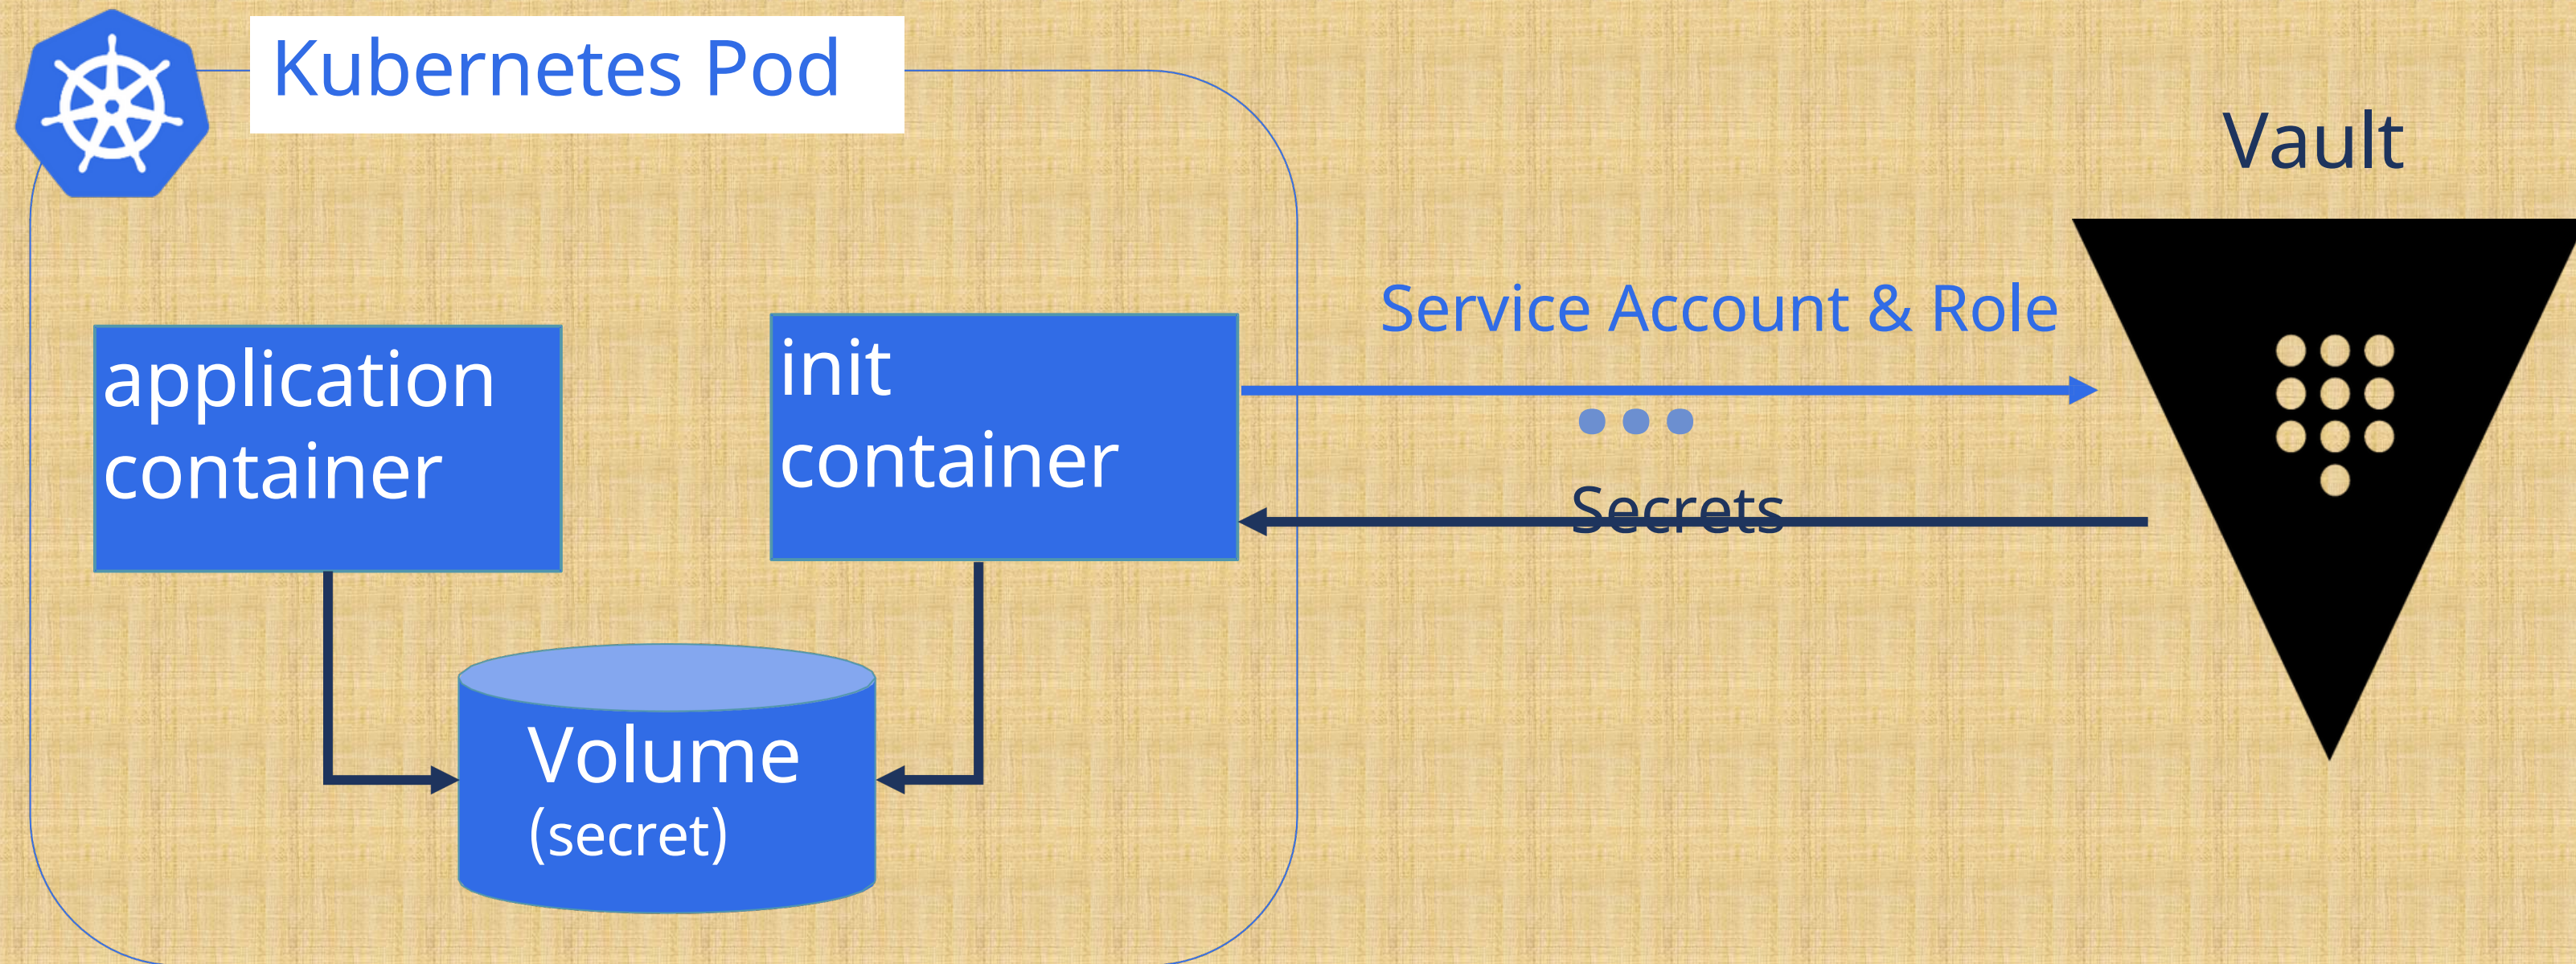


# Secrets Management with Hashicorp Vault

- Helps automates security best practices for
  - Secrets Management
  - Auditing
  - Certificate Management
  - Encryption
- Dynamic Secrets
  - Credentials (keys, passwords, certificates) are generated when a client requests them
  - Credentials are per client
  - Credentials are automatically deleted if a lease expires



# An init container to fetch secrets



# Namespaces

- **Kubernetes Data Plane Virtualization**

Kubernetes supports multiple virtual clusters backed by the same physical cluster. These virtual clusters are called namespaces.

- Namespaces partition the Kubernetes object model so multiple objects with the same name can exist in the same cluster
- Namespaces are the foundation for applying other security constructs

# Role-based access control (RBAC)

- Users are authenticated via OIDC, X.509 certificates, tokens, etc.
- The authentication result can provide user and group information.
- However, Users and User Groups are managed externally (e.g. in an LDAP / AD server).
- Kubernetes has a fine-grained permission model
  - Role (namespace) / ClusterRole
- Roles are mapped to users or groups via role bindings
  - RoleBinding (namespace) / ClusterRoleBinding



# Service Accounts

- Service Accounts are meant for authenticating and authorizing processes
- Each namespace has a default service account
- Each Pod has a service account (default – if not specified)
- A best practice is to use a service account per app
- To prevent a service account token from being mounted in a Pod use “**automountServiceAccountToken: false**”. This can be enforced via a policy.

# Network Segmentation via Network Policies

- By default, Kubernetes pods are **“non-isolated”**
  - They accept network connections from any source and can initiate connection requests to any destination
- Network Policies define traffic rules for Kubernetes pods
  - ingress (inbound traffic)
  - egress (outbound traffic)

## Network Policy

### Pod Selector

### Ingress



Ingress Rule

### Egress



Egress Rule

# Resource Management


- Pods can have resource requests and limits
- This allows three quality of service models



**Burstable**



**Guaranteed**



- A namespace can have limits and default allocations
- Quotas and limits ensure fairness and stability

# Pod Security Policies

- Controls runtime security settings for pods
- Enabled at the API Controller
- Requires a role binding between pod Service Account and the PSP

Control Aspect	Field Names
Running of privileged containers	<a href="#">privileged</a>
Usage of host namespaces	<a href="#">hostPID</a> , <a href="#">hostIPC</a>
Usage of host networking and ports	<a href="#">hostNetwork</a> , <a href="#">hostPorts</a>
Usage of volume types	<a href="#">volumes</a>
Usage of the host filesystem	<a href="#">allowedHostPaths</a>
White list of Flexvolume drivers	<a href="#">allowedFlexVolumes</a>
Allocating an FSGroup that owns the pod's volumes	<a href="#">fsGroup</a>
Requiring the use of a read only root file system	<a href="#">readOnlyRootFilesystem</a>
The user and group IDs of the container	<a href="#">runAsUser</a> , <a href="#">runAsGroup</a> , <a href="#">supplementalGroups</a>
Restricting escalation to root privileges	<a href="#">allowPrivilegeEscalation</a> , <a href="#">defaultAllowPrivilegeEscalation</a>
Linux capabilities	<a href="#">defaultAddCapabilities</a> , <a href="#">requiredDropCapabilities</a> , <a href="#">allowedCapabilities</a>
The SELinux context of the container	<a href="#">seLinux</a>
The Allowed Proc Mount types for the container	<a href="#">allowedProcMountTypes</a>
The AppArmor profile used by containers	<a href="#">annotations</a>
The seccomp profile used by containers	<a href="#">annotations</a>
The sysctl profile used by containers	<a href="#">forbiddenSysctls</a> , <a href="#">allowedUnsafeSysctls</a>



# Use a policy engine to audit and enforce

- Pod Security Policies are tricky to manage
  - Require a role binding to SA
  - Applied in alphabetical order
- Kyverno supports enforcement of the important PSP checks

```
kind: ClusterPolicy
metadata:
  name: validate-deny-runasrootuser
spec:
  validationFailureAction: "audit"
  rules:
  - name: deny-runasrootuser
    exclude:
      resources:
        namespaces:
        - kube-system
    match:
      resources:
        kinds:
        - Pod
    validate:
      message: "Root user is not allowed. Set runAsNonRoot to true."
      anyPattern:
      - spec:
          securityContext:
            runAsNonRoot: true
      - spec:
          containers:
          - name: "*"
            securityContext:
```



**Thank You**